



NYU


Center for
Data Science

Week 04.1: Distributed Storage

DS-GA 1004: Big Data

Last time: Map-Reduce

- Two functions: **mapper** and **reducer**
- **Mapper** consumes inputs, produces output:
(*key*, *value*)
- **Reducer** consumes a single *key* and list of *values*,
and produces *values*



Map-Reduce details...

- How is data **shared** over the cluster?
- How do we handle **node failure**?
- How can we **optimize for Map-Reduce** operations?

This week


1. **Distributed storage**
2. The Hadoop distributed file system (HDFS)

\$ `hadoop fs -command ...`

3. HDFS and Map-reduce


Start simple

- Imagine implementing Map-Reduce from scratch
 - ... with all data located on a file server
- Head node sends (**mapper**, **reducer**) code to each worker + **block of data**
- Workers send **output** back to head
 - **Mappers** → intermediate results
 - **Reducers** → final results



Start simple

- Imagine implementing Map-Reduce from scratch
 - ... with all data located on a file server
- Head node sends (**mapper**, **reducer**) code to each worker + **block of data**
- Workers send **output** back to head
 - **Mappers** → intermediate results
 - **Reducers** → final results




This will work, but it's inefficient!

Each job moves the entire data set over the network!

This is a **failure of locality**.


Localize all data?

- What if all data is replicated on all worker nodes?
- Head node sends (**mapper**, **reducer**) code to each worker
+ id of data block
- Workers send **output** back to head
 - **Mappers** → intermediate results
 - **Reducers** → final results



Localize all data?

- What if all data is replicated on all worker nodes?
- Head node sends (**mapper**, **reducer**) code to each worker
+ id of data block
- Workers send **output** back to head
 - **Mappers** → intermediate results
 - **Reducers** → final results



This also will work, but it's expensive!

Each worker needs a large amount of storage.

Most workers don't touch most of the data.

Design considerations


- Communication costs (bytes transferred)
- Fault tolerance
- Redundancy vs. communication
- Granularity of access
- Locality
- Common access patterns
- **Programs are small, data is big!**

Background: storage systems

Disks and disk arrays

Distributed file systems

- **Distributed file systems** store data over many machines
- Different from **networked file systems** (NFS), which make centrally hosted data transparently available on many machines
- First, some background on storage systems...
 - (This will be fast and irresponsibly high-level)



File systems and hard disks

- File systems are made of **directories** and **files**
- Disks are made of **contiguous sectors**
 - Typically 512 or 4096 bytes
 - This is the smallest addressable unit of storage
 - Each sector belongs to at most one file
- How are these reconciled?

```
bmcfee@mariana.cims.nyu.edu /scratch/bmcfee/data/MSD  
→ find data | head -30 | grep h5  
data/A/A/A/TRAAZF12903CCCF6B.h5  
data/A/A/A/TRAAAK128F9318786.h5  
data/A/A/A/TRAAAYX128F4263BC0.h5  
data/A/A/A/TRAAAV128F421A322.h5  
data/A/A/A/TRAAAW128F429D538.h5  
data/A/A/A/TRAAAY128F42A73F0.h5  
data/A/A/A/TRAAABD128F429CF47.h5  
data/A/A/A/TRAAACN128F9355673.h5  
data/A/A/A/TRAAACV128F423E09E.h5  
data/A/A/A/TRAADJ128F4287B47.h5  
data/A/A/A/TRAADT12903CCC339.h5
```



Files and blocks


- **Files** are broken into **blocks**

- Block size \geq sector size




- Each **block** is mapped onto **sectors**
- The file system (OS) hides this from us
- **Result:** a single file can spread over the entire disk

```
bmcfee@mariana.cims.nyu.edu /scratch/bmcfee/data/MSD
→ find data | head -30 | grep h5
data/A/A/A/TRAAAZF12903CCCF6B.h5
data/A/A/A/TRAAAK128F9318786.h5
data/A/A/A/TRAAAYX128F4263BC0.h5
data/A/A/A/TRAAAV128F421A322.h5
data/A/A/A/TRAAAW128F429D538.h5
data/A/A/A/TRAAAY128F42A73F0.h5
data/A/A/A/TRAAABD128F429CF47.h5
data/A/A/A/TRAAACN128F9355673.h5
data/A/A/A/TRAAACV128F423E09E.h5
data/A/A/A/TRAADJ128F4287B47.h5
data/A/A/A/TRAADT12903CCC339.h5
```



Large-scale storage

- What if our data is **too big** for a single disk?
- What happens if a **disk fails**?
- Throughput is limited by moving the **head**
 - ... what if we have multiple simultaneous read/write requests?




Redundant array of inexpensive disks (RAID)

- **RAID** systems distribute storage over multiple disks in a single machine
 - But look like a single volume to the OS
- Goals of **RAID**:
 - **Capacity**
 - **Reliability**
 - **Throughput**
- Comes in multiple “levels” with different reliability-capacity trade-offs



Commonly used RAID levels





Striping only

No fault-tolerance

Capacity scales linearly

Commonly used RAID levels



Striping only

No fault-tolerance


Capacity scales linearly

Full redundancy

(n-1) fault-tolerance

Capacity is constant


Commonly used RAID levels



Striping only

No fault-tolerance


Capacity scales linearly



Full redundancy

(n-1) fault-tolerance

Capacity is constant



Striping with distributed **parity**

$$A_p = A_1 \text{ XOR } A_2 \text{ XOR } \dots \text{ XOR } A_{n-1}$$


Can tolerate some failure

Capacity scales almost linearly

Requires $n \geq 3$ disks

Throughput with replication

- **Write** speed **decreases**: all data must be pushed to **all disks**
- **Read** speed can **increase**: blocks can be read in parallel from **any disk**
- *When is this trade-off worth it?*



Features of RAID

- **Striping blocks** over multiple drives increases **capacity and throughput**
 - **Not** reliability
- **Striping blocks** \Rightarrow a file can be larger than any single drive
- Adding **parity blocks** improves reliability
 - If a data block is corrupted or lost, it can be recovered from other blocks + parity
 - If a parity block is corrupted or lost, it can be recomputed from data blocks

Journaling

- Some file system operations are **non-atomic**
 - Creating a file: 1) allocate blocks and inodes, 2) create directory entry
- What if the system crashes mid-operation?
- **Solution:** use a **journal** to stage operations to be completed
- Crash recovery **plays back the journal**, does not need to check entire disk!
 - Rolls back partially completed operations
 - Clears completed operations

Quick recap so far

- **Combining multiple disks** increases capacity
- **Redundant storage** increases reliability and read-throughput
- **Parity blocks** provide error detection
- **Journaling** provides fast recovery from system failures

Why wasn't RAID enough?

- RAID improves capacity, fault-tolerance, and (read) throughput **on a single machine**
- What about distributed computation?
 - **Communication** over the network is a **bottleneck**
- *What are the common access patterns?*
 - Can we do **better** than both **fully localized** and **fully distributed**?

Up next...

- Part 2: Hadoop distributed filesystem
- Part 3: HDFS + MapReduce