



NYU

Center for  
Data Science

# Week 06.3: Column storage

DS-GA 1004: Big Data

Lab 2 ...

# Announcements


- Lab 3 (Spark) starts Thursday
- No class next week!

# Previously...

- Spark and RDDs
- Delayed computation  $\Rightarrow$  optimization

# This week

1. Column-oriented storage
2. Dremel & Parquet



**vs.**



# Row-oriented storage: CSV files?

- Imagine you have data stored as rows of text in the usual way

```
id, name, mass  
1, T.Rex, 8000  
2, Stegosaurus, 4000  
3, Ankylosaurus, 4000  
...
```

# Row-oriented storage: CSV files?

- Imagine you have data stored as rows of text in the usual way

```
id, name, mass  
1, T.Rex, 8000  
2, Stegosaurus, 4000  
3, Ankylosaurus, 4000  
...
```



```
id, name, mass\n1, T.Rex, 8000\n2, Stegosaurus, 4000\n3, Ankylosaurus, 4000\n...
```

# Row-oriented storage: CSV files?

- Imagine you have data stored as rows of text in the usual way

```
id, name, mass
1, T.Rex, 8000
2, Stegosaurus, 4000
3, Ankylosaurus, 4000
...
```



```
id, name, mass\n1, T.Rex, 8000\n2, Stegosaurus, 4000\n3, Ankylosaurus, 4000\n...
```

- How would you access the 1000th record?
- How would you access just the third column?



# Row-oriented storage: CSV files?

- Imagine you have data stored as rows of text in the usual way

```
id, name, mass
1, T.Rex, 8000
2, Stegosaurus, 4000
3, Ankylosaurus, 4000
...
```



```
id, name, mass\n1, T.Rex, 8000\n2, Stegosaurus, 4000\n3, Ankylosaurus, 4000\n...
```

## Problems:

- How would you access the 1000th record?
  - How would you access just the third column?
- Records are variable-length
  - Row and column offsets are hard to predict
  - Basically requires full serial scan

## Latency Numbers Every Programmer Should Know

■ 1ns  
■ L1 cache reference: 1ns  
■■■ Branch mispredict: 3ns  
■■■■ L2 cache reference: 4ns  
■■■■■ Mutex lock/unlock: 17ns

Send 2,000 bytes over commodity network: 44ns

Read 1,000,000 bytes sequentially from SSD: 49,000ns  $\approx$  49 $\mu$ s

SSD random read: 16,000ns  $\approx$  16 $\mu$ s

Disk seek: 2,000,000ns  $\approx$  2ms

Read 1,000,000 bytes sequentially from memory: 3,000ns  $\approx$  3 $\mu$ s

Read 1,000,000 bytes sequentially from disk: 825,000ns  $\approx$  825 $\mu$ s

- Transferring from disk to memory is incredibly slow
- Sequential memory reads are faster due to cache pre-fetching
- Strategies:
  - Transfer fewer bytes
  - Use predictable and contiguous memory access patterns

# Compression

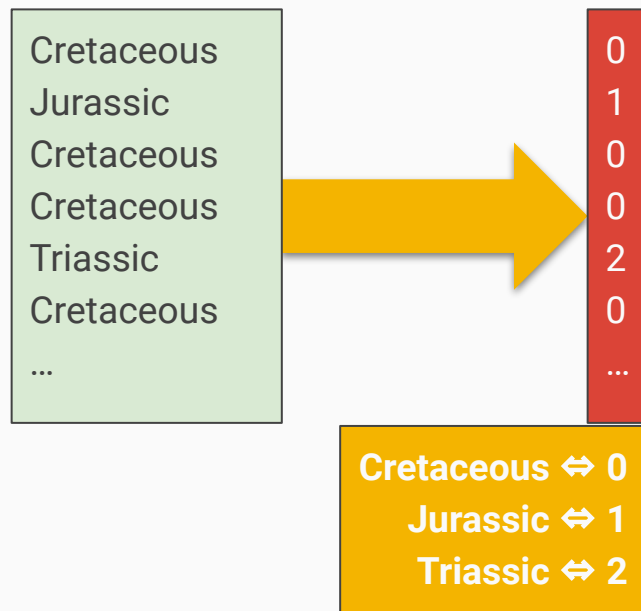
- **Records** have **heterogeneous** types
- A **single column** only has **one type**
- **Low entropy** in a column  $\Rightarrow$  **compression**
  - Compressed columns take less space
  - Compressed columns are **cheaper to load**
  - Sometimes we can **compute directly on compressed columns!**

id	Species	Era	Diet	Awesome	Mass
1	T. Rex	Cretaceous	Carnivore	True	8000
2	Stegosaurus	Jurassic	Herbivore	True	4000
3	Ankylosaurus	Cretaceous	Herbivore	False	4000

# Dictionary encoding

id	Species	Era	Diet	Awesome	Mass
1	T. Rex	Cretaceous	Carnivore	True	8000
2	Stegosaurus	Jurassic	Herbivore	True	4000
3	Ankylosaurus	Cretaceous	Herbivore	False	4000

- Useful when you have an attribute which takes **few distinct values**
- Replace **string values** by **string identifiers**
- Column now has **uniform data width**  
⇒ better cache locality!



slido



**When using dictionary coding, do we need to decompress the data to do partial matching?ex: "SELECT \* FROM Table WHERE name LIKE 'Sue%'"**

① Start presenting to display the poll results on this slide.

slido



**[2] When using dictionary coding, do we need to decompress the data to do partial matching?ex: "SELECT \* FROM Table WHERE name LIKE 'Sue%'"**

① Start presenting to display the poll results on this slide.

# Dictionary coding + partial matching

- We don't need to decompress!
- Do the partial match on the dictionary first
  - $\Rightarrow$  find all (if any) matching indices
- Then search the compressed table for the **matching indices**

# Bit-packing

- Integers usually consume 4, or 8 bytes (32 or 64 bits)
- Bit-packing** squeezes **small integers** together

Values	0	1	0	2	1	1
8-bit (binary)	0000 00 <u>00</u>	0000 00 <u>01</u>	0000 00 <u>00</u>	0000 00 <u>10</u>	0000 00 <u>01</u>	0000 00 <u>01</u>
Compressed	<u>0001</u> <u>0010</u>	<u>0101</u> ....				


- Matching and comparing can be done on **compressed values**



# Run-length encoding

id	Species	Era	Diet	Awesome	Mass
1	T. Rex	Cretaceous	Carnivore	True	8000
2	Stegosaurus	Jurassic	Herbivore	True	4000
3	Ankylosaurus	Cretaceous	Herbivore	False	4000

- Useful when you have long runs of a constant value
- Convert **sequence of values** to tuples **(value, # repetitions)**
- Sums, averages, counts, etc can all be done on **compressed values**



# Compression schemes abound...

- Frame of reference coding
  - 1004, 1005, 1006  $\Rightarrow$  **1000** | 4, 5, 6
- Delta coding
  - 1004, 1005, 1006  $\Rightarrow$  **1004** | +0, +1, +1
- Lempel-Ziv-Welch (LZW) compression

Compression schemes can be **combined!**

Delta + bit packing

Dictionary + Run-length encoding

Main trade-off is ***space efficiency*** vs.  
***complexity of querying/processing.***

slido



**Which of the following compression schemes are sensitive to the order of values within a column?**

① Start presenting to display the poll results on this slide.

slido



**[2] Which of the following compression schemes are sensitive to the order of values within a column?**

① Start presenting to display the poll results on this slide.

# Compression and ordering

- **Dictionary coding** doesn't depend on order
  - (Arguably, isn't even truly *compression*)
- **Bit-packing** depends only on the maximum value in a column
- **RLE, FoR, Delta** all depend critically on order!
  - Compare RLE on a,b,a,b,a,b,a,b vs. a,a,a,a,b,b,b,b
- **Take-away message: it can pay off to sort your data!**

# Nested and structured data

- Not everything fits nicely in relations / tables / dataframes
- Variable-length and variable-depth data can be difficult to deal with
- Record-oriented storage is relatively straightforward

**How can we get all the benefits of column stores but for structured data?**

# Record flattening

- **Key idea:** track repetitions of fields within a record
- **Repetition level (*r*):** which level repeated most recently?
- **Definition level (*d*):** how many optional fields in the path are present?
- **Required fields  $\Rightarrow$  Same levels as parent**
- **Optional fields  $\Rightarrow$  Same *r*-level as parent, *d*-level increments**
- **Repeated fields  $\Rightarrow$  *r*-level and *d*-level both increment from parent**



```
DocID: 10
Links:
    Forward: 20
    Forward: 40
    Forward: 60
Name:
    Language:
        Code: 'en-us'
        Country: 'us'
    Language:
        Code: 'en'
    URL: 'http://A'
Name:
    URL: 'http://B'
Name:
    Language:
        Code: 'en-gb'
        Country: 'gb'
```

# Partial record assembly

- Dremel can rebuild **partial views** (projections) of the data easily
- Unused attributes can be ignored!
- But decoding is **inherently sequential**  $\Rightarrow$  difficult to parallelize

Node.DocID		
value	r	d
10	0	1
20	0	1

Node.Links.Forward		
value	r	d
20	0	2
40	2	2
60	2	2
80	0	2

Node.Links.Backward		
value	r	d
NULL	0	1
10	0	2
30	1	2



DocID: 10  
Links:  
    Forward: 20  
    Forward: 40  
    Forward: 60

DocID: 20  
Links:  
    Backward: 10  
    Backward: 30  
    Forward: 80

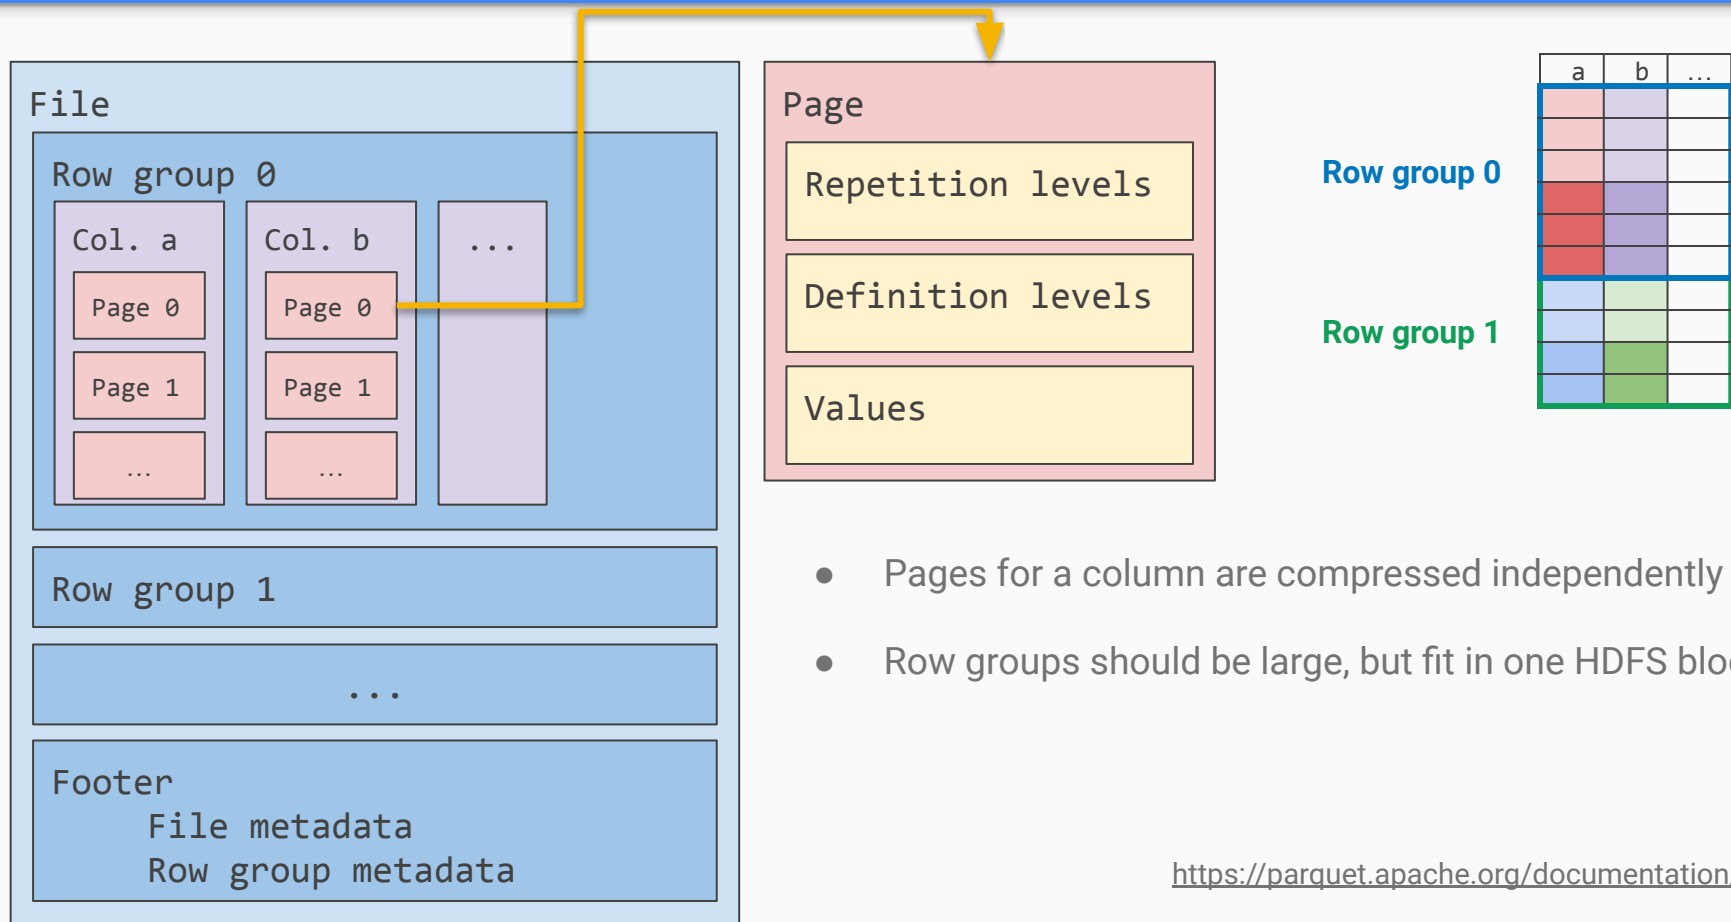


# After flattening...

- Repetition and definition columns are highly compressible
  - Not even needed for complete, tabular data!
- Value fields are now columnar
  - May also be compressed
- Columns are broken into **blocks** and compressed independently
  - This alleviates some decoding complexity and improves parallelism

value	r	d
http://A	0	2
http://B	1	2
NULL	1	1
http://C	0	2

# Parquet format



- Pages for a column are compressed independently
- Row groups should be large, but fit in one HDFS block

slido



# What are some benefits and drawbacks of using large pages?

① Start presenting to display the poll results on this slide.

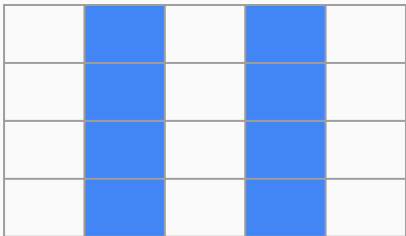
# Page size in parquet

- Larger pages
  - Potentially better compression rate
  - But worse overhead for serial decoding
- Smaller pages
  - Faster decoding: fewer records to scan through before reconstructing
  - Worse compression rate due to less context

# Column storage take-aways

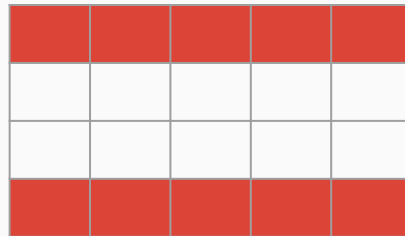
## Pros:

- Can be much faster when you only want a **subset of attributes**
- Higher **storage efficiency** and **throughput**
- Collecting **data of the same type** enables compression and better access patterns




## Cons:

- Reconstructing full tuples can be slow
  - Not great for **record-oriented** jobs
- Writes / deletion can be slow
- Handling non-tabular data is tricky
  - Dremel → Parquet to the rescue




# After spring break...

- Dask (last software/framework)
- Then on to other things:
  - Data structures
  - Algorithms
  - Applications

# Flattening example

**Node.DocID**

value	r	d
-------	---	---

**Node.Name.URL**

value	r	d
-------	---	---

**Node.Links.Forward**

value	r	d
-------	---	---

**Node.Links.Backward**

value	r	d
-------	---	---

**Node.Name.Language.Code**

value	r	d
-------	---	---

**Node.Name.Language.Country**

value	r	d
-------	---	---

DocID: 10

Links:

Forward: 20

Forward: 40

Forward: 60

Name:

Language:

Code: 'en-us'

Country: 'us'

Language:

Code: 'en'

URL: 'http://A'

Name:

URL: 'http://B'

Name:

Language:

Code: 'en-gb'

Country: 'gb'

DocID: 20

Links:

Backward: 10

Backward: 30

Forward: 80

Name:

URL: 'http://C'

# Flattening example

DocID: 10

Links:

Forward: 20  
Forward: 40  
Forward: 60

Name:

Language:  
Code: 'en-us'  
Country: 'us'

Language:  
Code: 'en'  
URL: 'http://A'

Name:

URL: 'http://B'

Name:

Language:  
Code: 'en-gb'  
Country: 'gb'

Node.DocID

value	r	d
10	0	0

Node.Name.URL

value	r	d
-------	---	---

Node.Links.Forward

value	r	d
-------	---	---

Node.Links.Backward

value	r	d
-------	---	---

Node.Name.Language.Code

value	r	d
-------	---	---

Node.Name.Language.Country

value	r	d
-------	---	---

DocID is required

r=0, d=0

DocID: 20

Links:

Backward: 10  
Backward: 30  
Forward: 80

Name:

URL: 'http://C'



# Flattening example

Node.DocID

value	r	d
10	0	0

Node.Name.URL

value	r	d
-------	---	---

Node.Links.Forward

value	r	d
-------	---	---

Node.Links.Backward

value	r	d
NULL	0	1

Node.Name.Language.Code

value	r	d
-------	---	---

Node.Name.Language.Country

value	r	d
-------	---	---

Links is **optional** (but present)  
Links.Backward is **repeated** (but absent)

r=0, d=1

No value in this record, so fill a NULL

DocID: 10

Links:

Forward: 20

Forward: 40

Forward: 60

Name:

Language:

Code: 'en-us'

Country: 'us'

Language:

Code: 'en'

URL: 'http://A'

Name:

URL: 'http://B'

Name:

Language:

Code: 'en-gb'

Country: 'gb'

DocID: 20

Links:

Backward: 10

Backward: 30

Forward: 80

Name:

URL: 'http://C'

# Flattening example

Node.DocID		
value	r	d
10	0	0

Node.Name.URL		
value	r	d

Node.Links.Forward		
value	r	d
20	0	2

Node.Links.Backward		
value	r	d
NULL	0	1

Node.Name.Language.Code		
value	r	d

Node.Name.Language.Country		
value	r	d

DocID: 10

Links:

**Forward: 20**

Forward: 40

Forward: 60

Name:

Language:

Code: 'en-us'

Country: 'us'

Language:

Code: 'en'

URL: 'http://A'

Name:

URL: 'http://B'

Name:

Language:

Code: 'en-gb'

Country: 'gb'

...Forward  $\Rightarrow$  d=2

No repetitions: r=0

DocID: 20

Links:

Backward: 10

Backward: 30

Forward: 80

Name:

URL: 'http://C'

# Flattening example

Node.DocID

value	r	d
10	0	0

Node.Name.URL

value	r	d
-------	---	---

Node.Links.Forward

value	r	d
20	0	2
40	1	2

Node.Links.Backward

value	r	d
NULL	0	1

Node.Name.Language.Code

value	r	d
-------	---	---

Node.Name.Language.Country

value	r	d
-------	---	---

DocID: 10

Links:

Forward: 20

**Forward: 40**

Forward: 60

Name:

Language:

Code: 'en-us'

Country: 'us'

Language:

Code: 'en'

URL: 'http://A'

Name:

URL: 'http://B'

Name:

Language:

Code: 'en-gb'

Country: 'gb'

...Forward  $\Rightarrow$  d=2

Repetition in level r=1

DocID: 20

Links:

Backward: 10

Backward: 30

Forward: 80

Name:

URL: 'http://C'

# Flattening example

Node.DocID

value	r	d
10	0	0

Node.Name.URL

value	r	d
-------	---	---

Node.Links.Forward

value	r	d
20	0	2
40	1	2
60	1	2

Node.Links.Backward

value	r	d
NULL	0	1

Node.Name.Language.Code

value	r	d
-------	---	---

Node.Name.Language.Country

value	r	d
-------	---	---

DocID: 10

Links:

Forward: 20

Forward: 40

**Forward: 60**

Name:

Language:

Code: 'en-us'

Country: 'us'

Language:

Code: 'en'

URL: 'http://A'

Name:

URL: 'http://B'

Name:

Language:

Code: 'en-gb'

Country: 'gb'

...Forward  $\Rightarrow$  d=2

Repetition in level r=1

DocID: 20

Links:

Backward: 10

Backward: 30

Forward: 80

Name:

URL: 'http://C'

# Flattening example

**Node.DocID**

value	r	d
10	0	0

**Node.Name.URL**

value	r	d
-------	---	---

**Node.Links.Forward**

value	r	d
20	0	2
40	1	2
60	1	2

**Node.Links.Backward**

value	r	d
NULL	0	1

**Node.Name.Language.Code**

value	r	d
en-us	0	2

**Node.Name.Language.Country**

value	r	d
-------	---	---

Name.Language.Code required

First occurrence (r=0)  
Full definition path (d=2)

DocID: 10

Links:

Forward: 20

Forward: 40

Forward: 60

Name:

Language:

Code: 'en-us'

Country: 'us'

Language:

Code: 'en'

URL: 'http://A'

Name:

URL: 'http://B'

Name:

Language:

Code: 'en-gb'

Country: 'gb'

DocID: 20

Links:

Backward: 10

Backward: 30

Forward: 80

Name:

URL: 'http://C'

# Flattening example

Node.DocID

value	r	d
10	0	0

Node.Name.URL

value	r	d
-------	---	---

Node.Links.Forward

value	r	d
20	0	2
40	1	2
60	1	2

Node.Links.Backward

value	r	d
NULL	0	1

Node.Name.Language.Code

value	r	d
en-us	0	2

Node.Name.Language.Country

value	r	d
us	0	3

...Country is optional  $\Rightarrow d=3$

First occurrence (r=0)

Full definition path (d=3)

DocID: 10

Links:

Forward: 20

Forward: 40

Forward: 60

Name:

Language:

Code: 'en-us'

Country: 'us'

Language:

Code: 'en'

URL: 'http://A'

Name:

URL: 'http://B'

Name:

Language:

Code: 'en-gb'

Country: 'gb'

DocID: 20

Links:

Backward: 10

Backward: 30

Forward: 80

Name:

URL: 'http://C'

# Flattening example

**Node.DocID**

value	r	d
10	0	0

**Node.Name.URL**

value	r	d
-------	---	---

**Node.Links.Forward**

value	r	d
20	0	2
40	1	2
60	1	2

**Node.Links.Backward**

value	r	d
NULL	0	1

**Node.Name.Language.Code**

value	r	d
en-us	0	2
en	2	2

**Node.Name.Language.Country**

value	r	d
us	0	3

...Code is required

Repetition at r=2  
(Name.Language)

DocID: 10

Links:

Forward: 20

Forward: 40

Forward: 60

Name:

Language:

Code: 'en-us'

Country: 'us'

Language:

Code: 'en'

URL: 'http://A'

Name:

URL: 'http://B'

Name:

Language:

Code: 'en-gb'

Country: 'gb'

DocID: 20

Links:

Backward: 10

Backward: 30

Forward: 80

Name:

URL: 'http://C'

# Flattening example

Node.DocID

value	r	d
10	0	0

Node.Name.URL

value	r	d
-------	---	---

Node.Links.Forward

value	r	d
20	0	2
40	1	2
60	1	2

Node.Links.Backward

value	r	d
NULL	0	1

Node.Name.Language.Code

value	r	d
en-us	0	2
en	2	2

Node.Name.Language.Country

value	r	d
us	0	3
NULL	2	2

...Language.Country optional

Repeated at Language level  
r=2, d=2

DocID: 10

Links:

Forward: 20

Forward: 40

Forward: 60

Name:

Language:

Code: 'en-us'

Country: 'us'

Language:

Code: 'en'

URL: 'http://A'

Name:

URL: 'http://B'

Name:

Language:

Code: 'en-gb'

Country: 'gb'

DocID: 20

Links:

Backward: 10

Backward: 30

Forward: 80

Name:

URL: 'http://C'



# Flattening example

**Node.DocID**

value	r	d
10	0	0

**Node.Name.URL**

value	r	d
http://A	0	2

**Node.Links.Forward**

value	r	d
20	0	2
40	1	2
60	1	2

**Node.Links.Backward**

value	r	d
NULL	0	1

**Node.Name.Language.Code**

value	r	d
en-us	0	2
en	2	2

**Node.Name.Language.Country**

value	r	d
us	0	3
NULL	2	2

Node.Name.URL is optional  
⇒ d=2

No repetitions: r=0

DocID: 10

Links:

Forward: 20

Forward: 40

Forward: 60

Name:

Language:

Code: 'en-us'

Country: 'us'

Language:

Code: 'en'

URL: 'http://A'

Name:

URL: 'http://B'

Name:

Language:

Code: 'en-gb'

Country: 'gb'

DocID: 20

Links:

Backward: 10

Backward: 30

Forward: 80

Name:

URL: 'http://C'

# Flattening example

**Node.DocID**

value	r	d
10	0	0

**Node.Name.URL**

value	r	d
http://A	0	2

**Node.Links.Forward**

value	r	d
20	0	2
40		
60		

**Node.Links.Backward**

value	r	d
NULL	0	1

Node.Name ⇒ d=1

But no Language.\* data...

**Node.Name.Language.Code**

value	r	d
en-us	0	2
en	2	2
NULL	1	1

**Node.Name.Language.Country**

value	r	d
us	0	3
NULL	2	2
NULL	1	1

DocID: 10

Links:

Forward: 20

Forward: 40

Forward: 60

Name:

Language:

Code: 'en-us'

Country: 'us'

Language:

Code: 'en'

URL: 'http://A'

**Name:**

URL: 'http://B'

Name:

Language:

Code: 'en-gb'

Country: 'gb'

DocID: 20

Links:

Backward: 10

Backward: 30

Forward: 80

Name:

URL: 'http://C'

# Flattening example

**Node.DocID**

value	r	d
10	0	0

**Node.Name.URL**

value	r	d
http://A	0	2
http://B	1	2

**Node.Links.Forward**

value	r	d
20	0	2
40	1	2
60	1	2

**Node.Links.Backward**

value	r	d
NULL	0	1

**Node.Name.Language.Code**

value	r	d
en-us	0	2
en	2	2
NULL	1	1

**Node.Name.Language.Country**

value	r	d
us	0	3
NULL	2	2
NU		

Node.Name.URL  $\Rightarrow$  d=2

Repetition at r=1 (Node.Name)

DocID: 10

Links:

Forward: 20

Forward: 40

Forward: 60

Name:

Language:

Code: 'en-us'

Country: 'us'

Language:

Code: 'en'

URL: 'http://A'

Name:

URL: 'http://B'

Name:

Language:

Code: 'en-gb'

Country: 'gb'

DocID: 20

Links:

Backward: 10

Backward: 30

Forward: 80

Name:

URL: 'http://C'

# Flattening example

Node.DocID

value	r	d
10	0	0

Node.Name.URL

value	r	d
http://A	0	2
http://B	1	2

Node.Links.Forward

value	r	d
20	0	2
40	1	2
60	1	2

Node.Links.Backward

value	r	d
NULL	0	1

Node.Name.Language.Code

value	r	d
en-us	0	2
en	2	2
NULL	1	1
en-gb	1	2

Node.Name.Language.Country

value	r	d
us	0	3
NULL	2	2
NU		

...Language.Code ⇒ d=2

Repetition at r=1 (Node.Name)

DocID: 10

Links:

Forward: 20

Forward: 40

Forward: 60

Name:

Language:

Code: 'en-us'

Country: 'us'

Language:

Code: 'en'

URL: 'http://A'

Name:

URL: 'http://B'

Name:

Language:

Code: 'en-gb'

Country: 'gb'

DocID: 20

Links:

Backward: 10

Backward: 30

Forward: 80

Name:

URL: 'http://C'

# Flattening example

Node.DocID

value	r	d
10	0	0

Node.Name.URL

value	r	d
http://A	0	2
http://B	1	2

Node.Links.Forward

value	r	d
20	0	2
40		
60		

Node.Links.Backward

value	r	d
NULL	0	1

...Language.Country  $\Rightarrow$  d=3

Repetition at r=1 (Node.Name)

Node.Name.Language.Code

value	r	d
en-us	0	2
en	2	2
NULL	1	1
en-gb	1	2

Node.Name.Language.Country

value	r	d
us	0	3
NULL	2	2
NULL	1	1
gb	1	3

DocID: 10

Links:

Forward: 20

Forward: 40

Forward: 60

Name:

Language:

Code: 'en-us'

Country: 'us'

Language:

Code: 'en'

URL: 'http://A'

Name:

URL: 'http://B'

Name:

Language:

Code: 'en-gb'

Country: 'gb'

DocID: 20

Links:

Backward: 10

Backward: 30

Forward: 80

Name:

URL: 'http://C'

# Flattening example

Node.DocID

value	r	d
10	0	0

Node.Name.URL

value	r	d
http://A	0	2
http://B	1	2
NULL	1	1

Node.Links.Forward

value	r	d
20	0	2
40	1	2
60	1	2

Node.Links.Backward

value	r	d
NULL	0	1

Node.Name.Language.Code

value	r	d
en-us	0	2
en	2	2
NULL	1	1
en-gb	1	2

Node.Name.Language.Country

value	r	d
us	0	3
NULL	2	2
NU		

Node.Name  $\Rightarrow$  d=1

No URL data

DocID: 10

Links:

Forward: 20

Forward: 40

Forward: 60

Name:

Language:

Code: 'en-us'

Country: 'us'

Language:

Code: 'en'

URL: 'http://A'

Name:

URL: 'http://B'

**Name:**

**Language:**

**Code: 'en-gb'**

**Country: 'gb'**

DocID: 20

Links:

Backward: 10

Backward: 30

Forward: 80

Name:

URL: 'http://C'

# Flattening example

**Node.DocID**

value	r	d
10	0	0
20	0	0

**Node.Name.URL**

value	r	d
http://A	0	2
http://B	1	2
NULL	1	1

**Node.Links.Forward**

value	r	d
20	0	2
40	1	2
60	1	2

**Node.Links.Backward**

value	r	d
NULL	0	1

**Node.Name.Language.Code**

value	r	d
en-us	0	2
en	2	2
NULL	1	1
en-gb	1	2

**Node.Name.Language.Country**

value	r	d
us	0	3
NULL	2	2
NU		

Node.DocID  $\Rightarrow$  d=0

Required field, new document  
(r=0)

DocID: 10

Links:

Forward: 20

Forward: 40

Forward: 60

Name:

Language:

Code: 'en-us'

Country: 'us'

Language:

Code: 'en'

URL: 'http://A'

Name:

URL: 'http://B'

Name:

Language:

Code: 'en-gb'

Country: 'gb'

DocID: 20

Links:

Backward: 10

Backward: 30

Forward: 80

Name:

URL: 'http://C'

# Flattening example

**Node.DocID**

value	r	d
10	0	0
20	0	0

**Node.Name.URL**

value	r	d
http://A	0	2
http://B	1	2
NULL	1	1

**Node.Links.Forward**

value	r	d
20	0	2
40	1	2
60	1	2

**Node.Links.Backward**

value	r	d
NULL	0	1
10	0	2

**Node.Name.Language.Code**

value	r	d
en-us	0	2
en	2	2
NULL	1	1
en-gb	1	2

**Node.Name.Language.Country**

value	r	d
us	0	3
NULL	2	2
NU		

Node.Links.Backward ⇒ d=2

DocID: 10

Links:

Forward: 20

Forward: 40

Forward: 60

Name:

Language:

Code: 'en-us'

Country: 'us'

Language:

Code: 'en'

URL: 'http://A'

Name:

URL: 'http://B'

Name:

Language:

Code: 'en-gb'

Country: 'gb'

DocID: 20

Links:

**Backward: 10**

Backward: 30

Forward: 80

Name:

URL: 'http://C'



# Flattening example

**Node.DocID**

value	r	d
10	0	0
20	0	0

**Node.Name.URL**

value	r	d
http://A	0	2
http://B	1	2
NULL	1	1
http://C	0	2

**Node.Links.Forward**

value	r	d
20	0	2
40	1	2
60	1	2
80	0	2

**Node.Links.Backward**

value	r	d
NULL	0	1
10	0	2
30	1	2

**Node.Name.Language.Code**

value	r	d
en-us	0	2
en	2	2
NULL	1	1
en-gb	1	2
NULL	0	1

**Node.Name.Language.Country**

value	r	d
us	0	3
NULL	2	2
NULL	1	1
gb	1	3
NULL	0	1

... and all  
the rest

DocID: 10

Links:

Forward: 20

Forward: 40

Forward: 60

Name:

Language:

Code: 'en-us'

Country: 'us'

Language:

Code: 'en'

URL: 'http://A'

Name:

URL: 'http://B'

Name:

Language:

Code: 'en-gb'

Country: 'gb'

DocID: 20

Links:

Backward: 10

Backward: 30

Forward: 80

Name:

URL: 'http://C'