



NYU

Center for
Data Science

Week 07.2: Dask / HPC

DS-GA 1004: Big Data

slido



We're half way through the semester. How are you doing?

① Start presenting to display the poll results on this slide.

Announcements

- Quiz 3 this week
 - Spark + Column-oriented storage
- Lab 3 due next week 04/01
- Lab 4 coming soon
 - Dask on the Greene cluster

This week

- Dask
- High-performance computing (HPC)
- Open Q&A

Roadmap for the semester

- | | | | | | |
|----|-------|------------------------------------|-----|-------|----------------------------|
| 1. | 01/24 | Introduction | 9. | 03/21 | HPC and Dask |
| 2. | 01/31 | Relational databases | 10. | 03/28 | Text and similarity search |
| 3. | 02/07 | Map-reduce | 11. | 04/04 | Reproducibility |
| 4. | 02/14 | Hadoop distributed file system | 12. | 04/11 | Recommender systems |
| 5. | 02/21 | <i>President's day, no meeting</i> | 13. | 04/18 | Graph algorithms |
| 6. | 02/28 | Spark | 14. | 04/25 | Differential privacy |
| 7. | 03/07 | Column-oriented storage | 15. | 05/02 | Graphical processing units |
| 8. | 03/14 | <i>Spring break, no meeting</i> | 16. | 05/09 | TBA |

Dask

[Rocklin, 2015]

- Python-based distributed computation
- Many common design principles with Spark
 - Delayed computation
 - Computation graphs
 - Collections-based interfaces (e.g. DataFrames)
- Some key differences:
 - Prioritizes array-based computation
 - Designed to support single-machine, out-of-core use

Delayed computation and task graphs

- Dask builds complex computations by composing deferred computations into a **task graph**


```
import dask
```

```
def square(x):  
    return x**2
```

```
f = dask.delayed(square)
```

```
y = f(5)
```


```
y.visualize() # draw computation graph
```



```
g = dask.delayed(sum)
```

```
z = g([f(x) for x in range(5)])
```

```
z.visualize()
```



Collections interfaces: Bags

- Dask **bags** are loosely analogous to Spark RDDs
- Unordered collection of generic Python objects
 - Partitions into subsets (sub-bags)
- Implements some basic operations
 - map, filter, join, sum, etc.
- A good choice for initial processing and structured objects
 - If your data is tabular or array-based, probably not the best choice

```
import dask.bag as db

b = db.from_sequence(range(5))

c = b.map(square)

c.compute() # [0, 1, 4, 9, 16]

c.sum().compute() # 30
```


Dask Bags vs. Spark RDDs

- Both partition a collection across multiple machines
- Both are immutable
- Both can be **transformed** Bag→Bag, RDD→RDD
- RDDs have **types** (e.g. `RDD[Integer]`)
- Bags are, more or less, **untyped**. You need to be more careful using them!

Bag folding vs grouping

- Try to **avoid using groupBy** on bags
 - This requires much inter-worker communication, and is slow!
- Use **foldBy** if possible
 - Similar benefits to a combiner in map-reduce
 - Perform local aggregation first to reduce the amount of shuffling
 - Like combiners, not always applicable, may require some cleverness
- You supply a **key function** and a **binary operation**

```
import dask.bag as db
```

```
b = db.from_sequence(range(10))
```

```
iseven = lambda x: x % 2 == 0
```


```
add = lambda x, y: x + y
```

```
dict(b.foldby(iseven, add))
```

Or better yet, move into **DataFrames** before any heavy processing...

Collections interfaces: DataFrames

- Just like you'd expect, similar to Spark DataFrames
 - Uses Pandas internally, interface is basically the same
- Parallelism (partitioning) is over subsets of **rows**
- Good choice for data that can naturally split into multiple CSV files (or Parquet partitions)




```
import dask.dataframe as dd
```

```
df = dd.read_csv('*.csv')  
df.mean().compute()
```

Collections interfaces: Arrays

- Dask Arrays work like NumPy arrays
- **Parallelism is not limited to rows**
 - You can define **chunks** along each dimension
- Large arrays are assembled implicitly from many small arrays
- Most* numpy operations work automatically




Array chunking example

```
import numpy as np  
import dask.array as da
```

```
# Make some noise: 2000 x 6000 x 5  
x = np.random.randn(2000, 6000, 5)
```

```
# Slice it up into chunks  
x_parallel = da.from_array(x, chunks=(1000, 1000, 2))
```



Things that don't work well with chunks

- Sorting. But you can often get away with **topk** instead!
- Operations where output size depends on values in data (e.g., masking: `x[x>0]`)
- Some linear algebra operations (**np.linalg**)

Handling large numerical data

- CSV and Parquet are not great options here!
- Collections of npy / npz files can be okay
 - `x.to_npy_stack('/output/folder/')` # Size and chunking info stored as metadata 'info' file
 - `y = da.from_npy_stack('/folder/containing/files/')`
- Hierarchical Data Format (**HDF5**) is a pretty good solution
 - Data can be accessed without fully loading into memory
- **ZARR** is a newer project with better distributed storage support

HDF5 (not the same as HDFS!)

- Basically a file-system within a file
 - (Hierarchical) Directory structures
- In python, use the **h5py** package
- Data is memory-mapped, not loaded
- Parallel reads are okay!

```
import h5py

data = h5py.File('myfile.h5', mode='r')

x = data['/x']
y = data['/y']
z = data['/path/to/z']
```

```
import dask.array as da

x_parallel = da.from_array(x, chunks=(1000, 1000))
```


Does Dask replace Spark?

- Eh... it depends 🙄
 - <https://docs.dask.org/en/latest/spark.html> summarizes use-cases and differences
- Pros for Dask:
 - Do you need to integrate with the SciPy stack? (Matplotlib, sklearn, etc)
 - Do you need to work with dense / multi-dimensional data?
 - Custom algorithms / advanced machine learning? GPUs?
- Pros for Spark:
 - More mature, possibly more stable / safe
 - More “high-level” -- you don’t need to think as much about the compute graph
 - Probably faster / better optimized for DataFrame crunching
 - Better support for large graph data

HPC: Peel and Greene

- So far, we've been using a Hadoop cluster (**Peel**)
 - HDFS storage
 - MapReduce + Spark jobs (YARN)
- We also have the **Greene** cluster for less restrictive computation
 - Network-accessible file storage (IBM GPFS, **not** HDFS)
 - SLURM job system: job scheduling is unaware of storage layout!
 - Traditionally preferred if you have “embarrassing parallelism” (i.e., independent computation / almost no communication)

Next week

- Approximate nearest neighbors
- Spatial data structures

slido



Audience Q&A Session

① Start presenting to display the audience questions on this slide.