



NYU

Center for  
Data Science

# Week 13.1: Graphics processing units

DS-GA 1004: Big Data

# This week

- **Graphics processing units (GPUs)**
- GPGPUs and CUDA
- Software frameworks

# Example: gradient descent (serial)

$$\min_w \sum_n f(x_n; w)$$

- Initialize  $w$
- **for**  $i \leftarrow 1$  **to** ITERATIONS:
  - Initialize  $G \leftarrow 0$
  - **for**  $n = 1 \dots N$ :
    - $G \mathrel{+}= \nabla_w f(x_n; w)$
  - $w \leftarrow w - G$

Total time:

ITERATIONS \*  $N$  \* [per-point gradient cost]

# Example: gradient descent (Spark version)

```
val points = spark.textFile(...).map(parsePoint).cache()
var w = Vector.random(D)

for (i ← 1 to ITERATIONS)
  val grad = spark.accumulator(new Vector(D))

  for (p ← points)
    val grad_p =  $\nabla_w f(p; w)$ 
    grad += grad_p

  w -= grad.value
```

Total time:


ITERATIONS \* ( $\sim N/k$ ) \* [per-point gradient cost]

- + Cost of communicating to **accumulator**
- + Cost of distributing data (**points**) and parameters (**w**)

# Alternative strategies

1. Gradient descent → Stochastic gradient descent (SGD)
  - a. Iterate on small batches instead of the whole dataset
  - b. In expectation, does the right thing (estimating gradient)
  - c. **Probably what you should be doing anyway!**
  - d. **But... requires a lot of communication!**
    - i. **Every iteration / minibatch uses different data**
    - ii. **weight vector is always changing**
2. **Use a different kind of computer**

Theoretical GFLOP/s at boost clock



# Why GPUs?


- Current best practice is to use GPUs to train deep networks
- **WHY? What about GPUs makes this work?**
- Can other processes be similarly accelerated by using GPUs?
- Speedups from parallelism usually come from constraints...
  - *What can't we accelerate with GPUs?*

# GPUs and the rendering pipeline



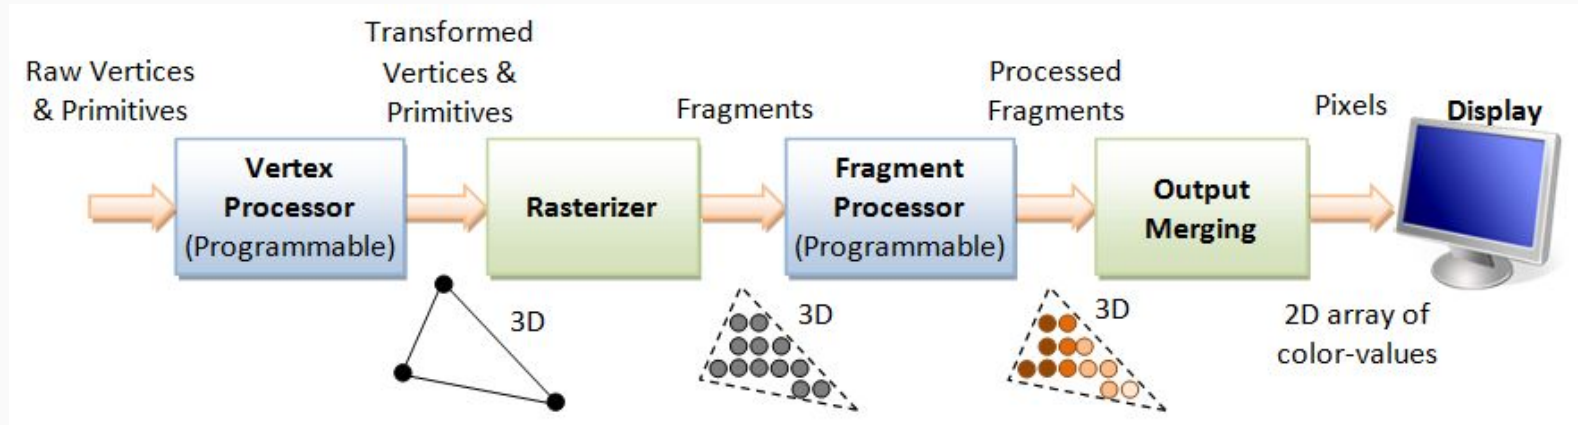
# 3D rendering

- Inputs:
  - 3D meshes
  - Textures, surface properties
  - Light sources
  - Camera position
- Outputs:
  - 2D array of pixels (rendered scene)
- Video games have real-time constraints



## Computational challenges:

- Scene complexity (# surfaces)
- Output resolution (# pixels)






## 1. Vertex processing


- Coordinate transformations for each model
- Camera transformation
- Camera lens / field of view / etc

Mostly linear or affine transformations  
(Rotation / translation / scaling)

### Outputs:

All vertices mapped to camera coordinates  
( $x, y, z$ )





## 2. Rasterizer


- Scans the scene for data to render at each pixel coordinate  $(x, y)$
- Object vertices don't necessarily line up to pixel coordinates  $\Rightarrow$  meshes are interpolated

### Outputs:

"Pixels" (or fragments) containing data to render at each pixel

### Notes:

Vertices  $\rightarrow$  Fragments is not generally 1-to-1  
This step is not programmable



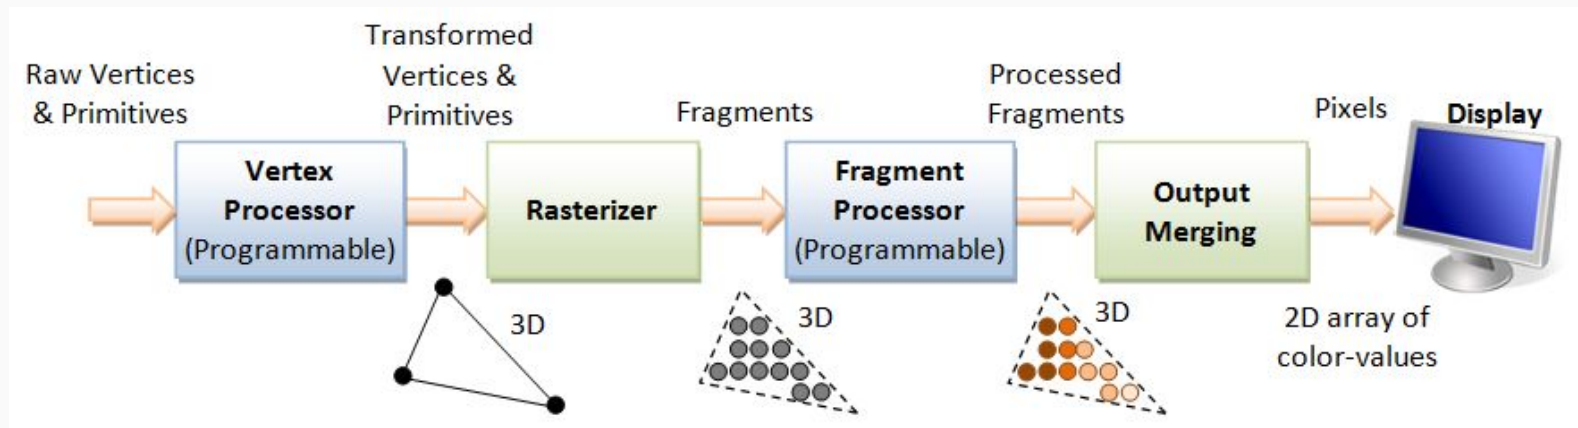
### 3. Fragment processing (aka “pixel shading”)

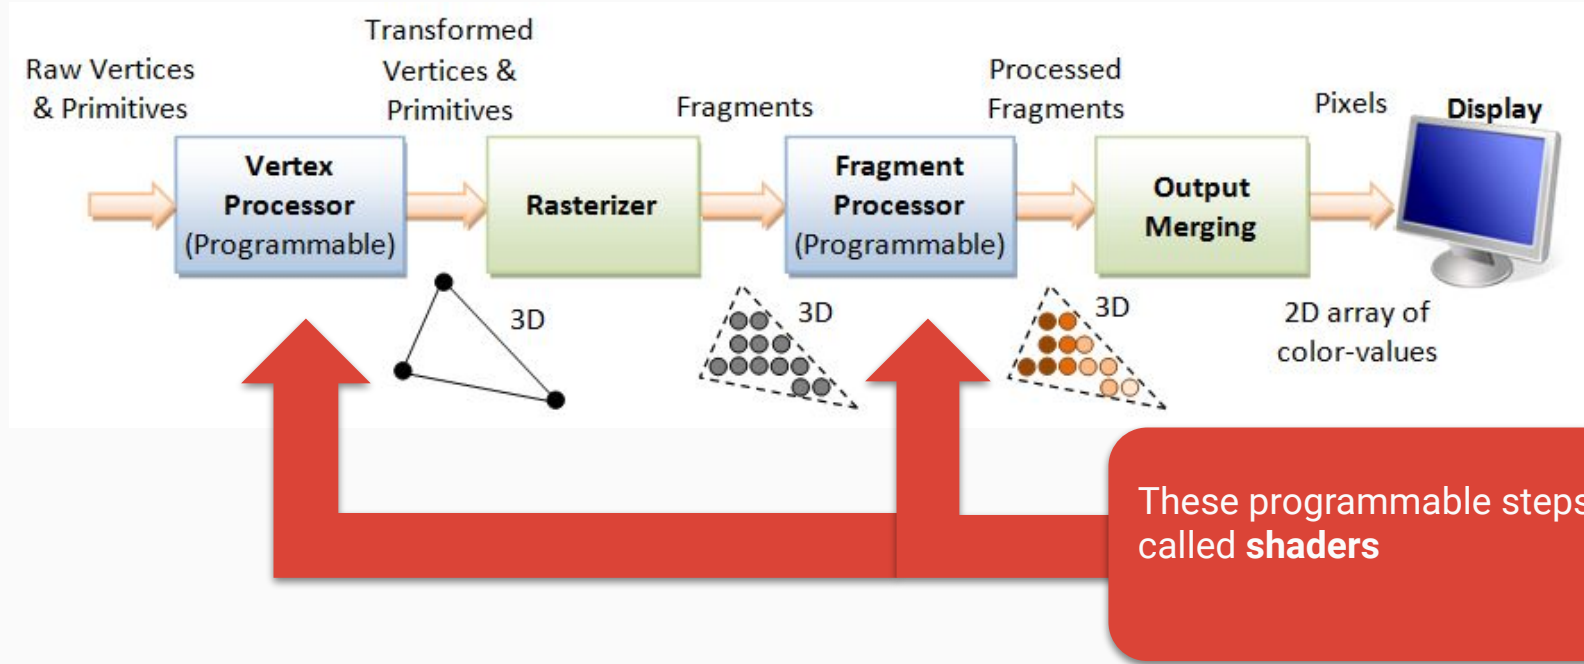

- Texture mapping, lighting
- Other visual effects (e.g. blur, masking, etc)

#### **Outputs:**

Color values for each fragment

May include occluded objects (discarded later)





# Parallelism in graphics

- Linear transformations can be applied **independently** to each **vertex**
- Texturing and lighting are also **independent** across **fragments**
- Specialized hardware can **parallelize** to meet real-time constraints
- To be cost-effective, each vertex or pixel processor needs to be **simple**




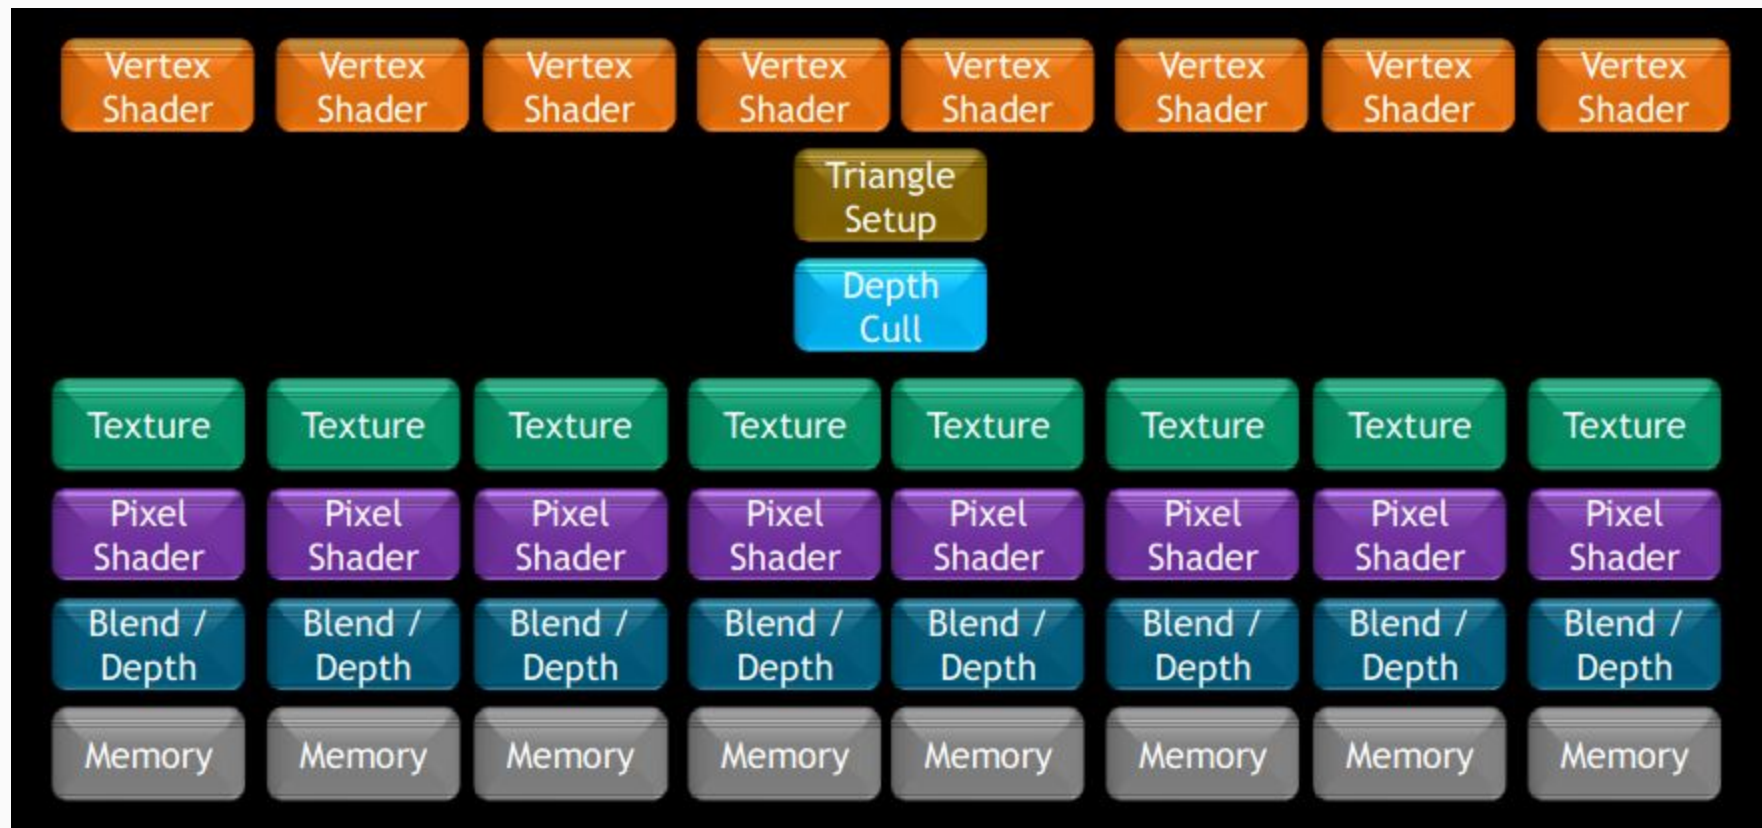
# Shaders

- “**Shaders**” are short programs that are applied independently to each vertex or fragment.

## Examples:

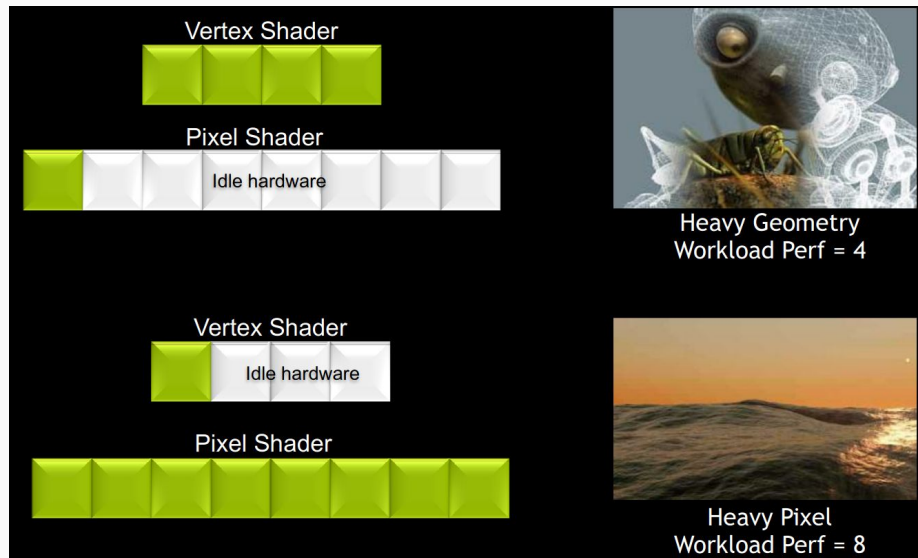
- Apply rotations to all vertices in a mesh
  - Set color value for a fragment by indexing a texture map
- Sounds a bit like a **mapper**, right?
- Shader code tends to have simple control flow





# Specialized shader units

- Older GPUs had separate processors for **vertex shading** and **pixel shading**
- This works well when the load is balanced, but real scenes rarely are!
  - **# Vertices**  $\neq$  **# Fragments**
- Unbalanced load means idle processors!
- Remember **key skew**?



# Summary

## Part 1: Graphics processing

- GPUs were designed to optimize for low-latency, highly parallel operations
- Shader programs are simple, computed in parallel, and combined
- Specialized shader units can suffer from imbalance similar to key-skew...
- Can we do something more general?

*... come back for part 2!*