



NYU

Center for
Data Science

Week 13.2:

General purpose GPUs

DS-GA 1004: Big Data

This week

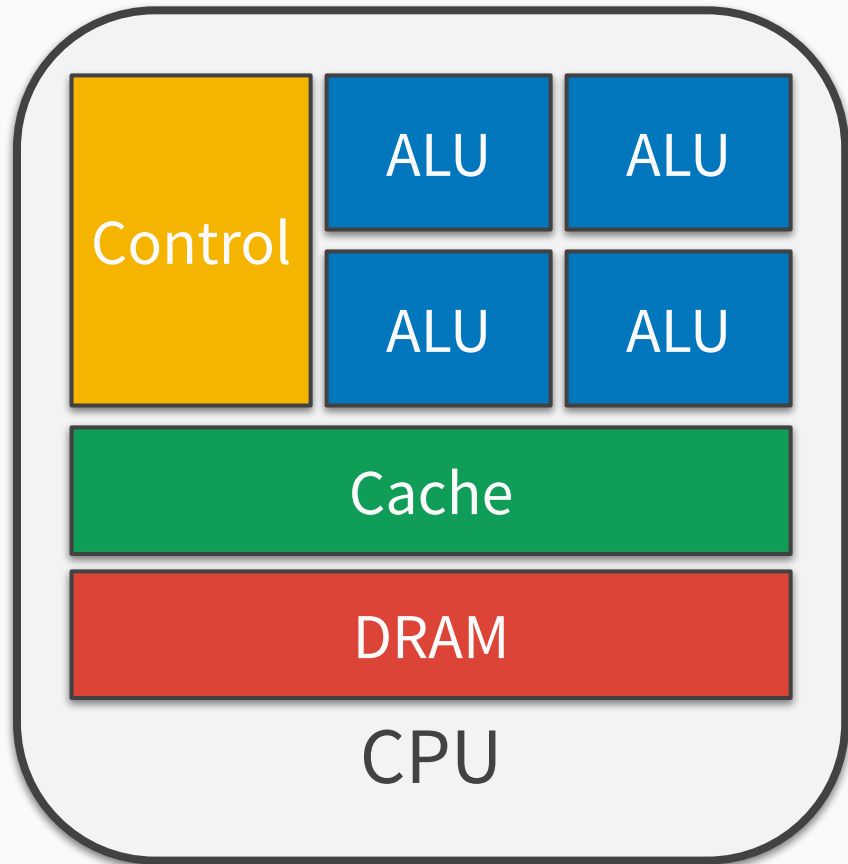
- Graphics processing units (GPUs)
- **GPGPUs and CUDA**
- Software frameworks

Shaders \Rightarrow Kernels

- General purpose GPUs (GPGPUs) remove the distinction between vertex- and pixel-cores
- “**Shader**” is replaced by “**kernel**” abstraction
 - This allows all cores to operate as either “vertex” or “pixel” cores...
 - Or as something else entirely!
- This is what NVidia’s **CUDA** (Compute Unified Device Architecture, 2006) API does

CPU vs GPU architecture

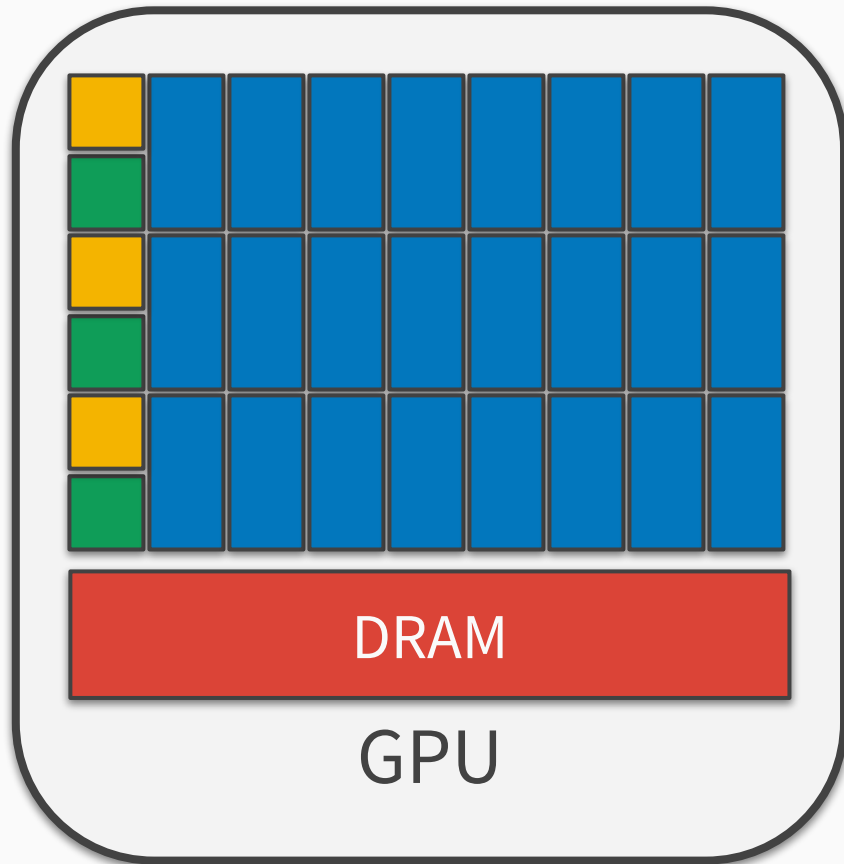
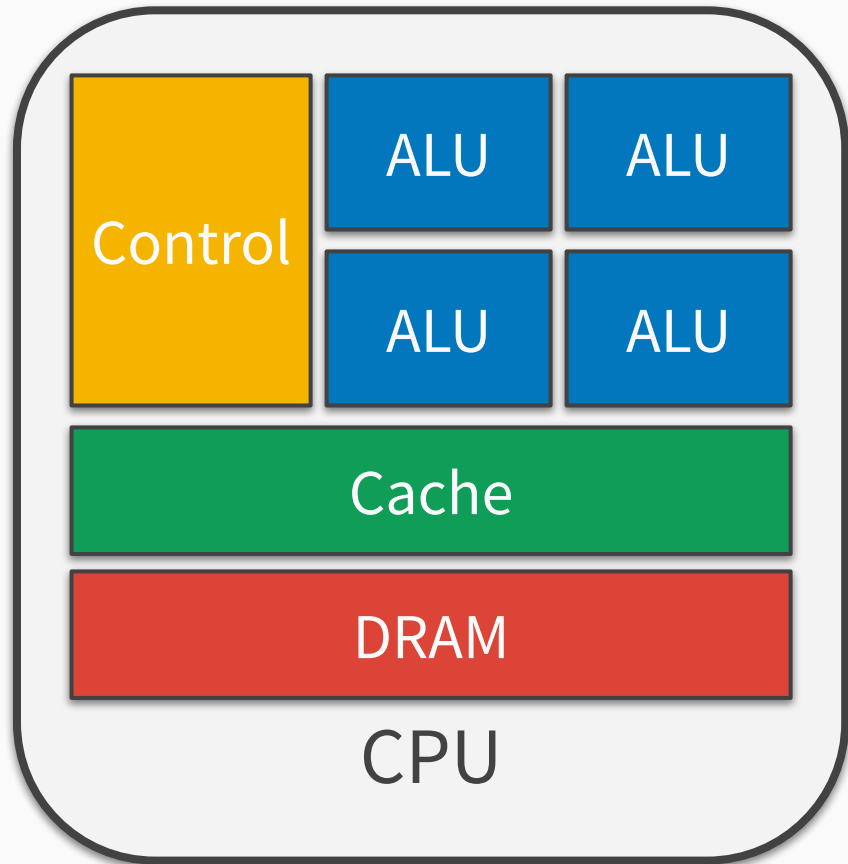
Adapted from <https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html>



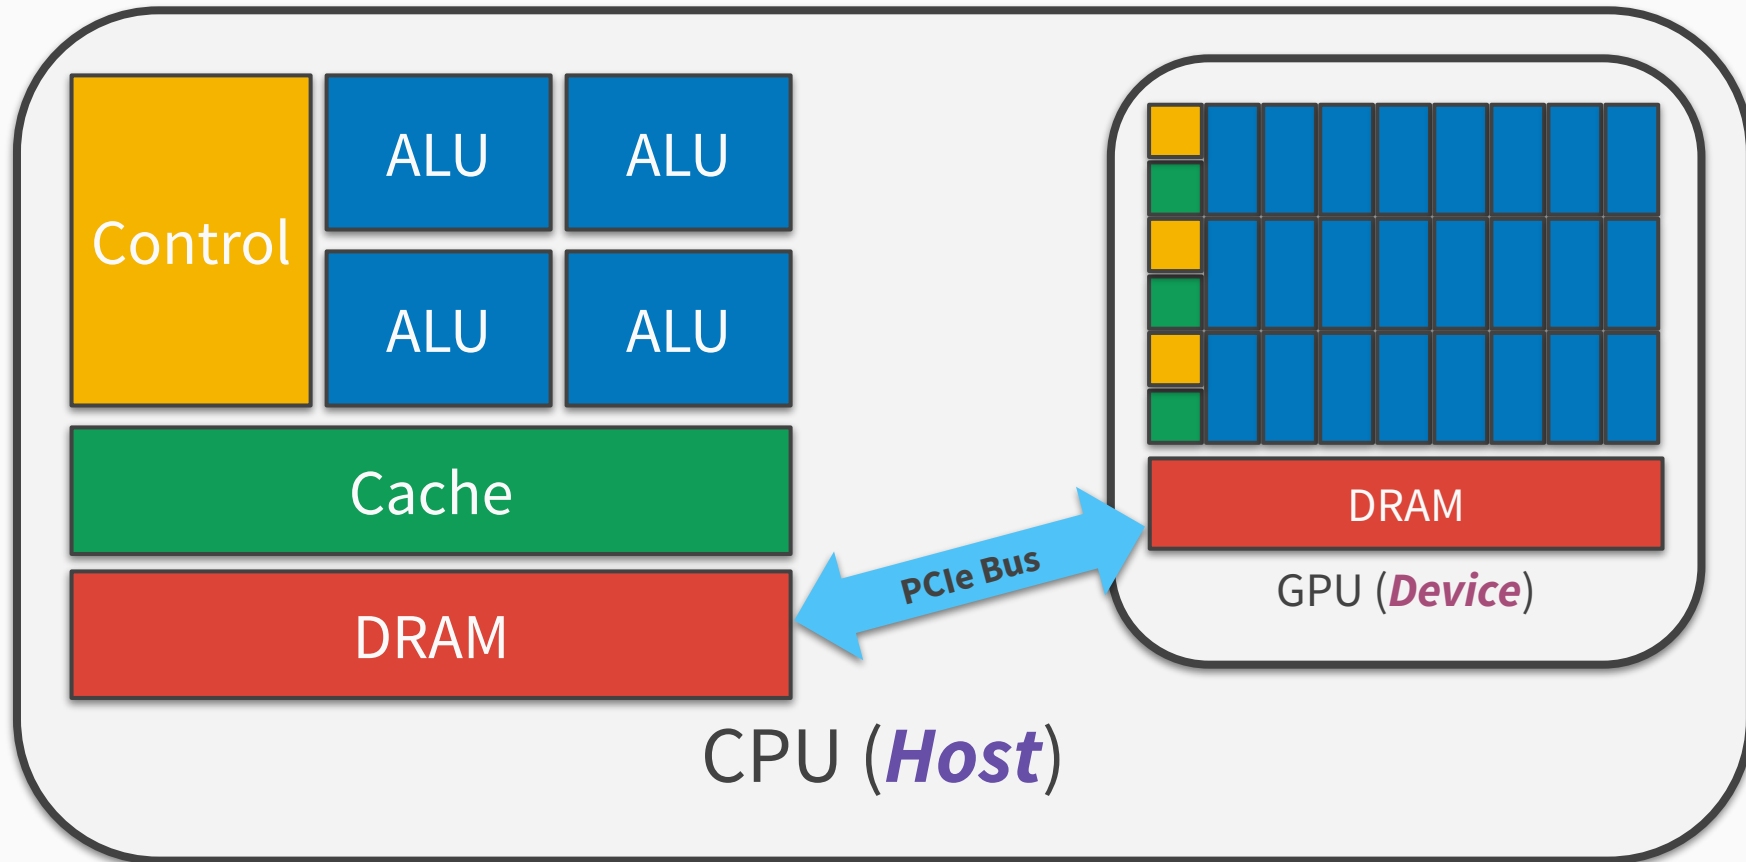
- **ALU** = Arithmetic logic unit (basic math operations)
- **Control** = program flow (branching, jumping, etc)
- **Cache** = on-CPU memory cache
- **DRAM** = Main system memory

CPU vs GPU architecture

Adapted from <https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html>



A slightly more accurate picture

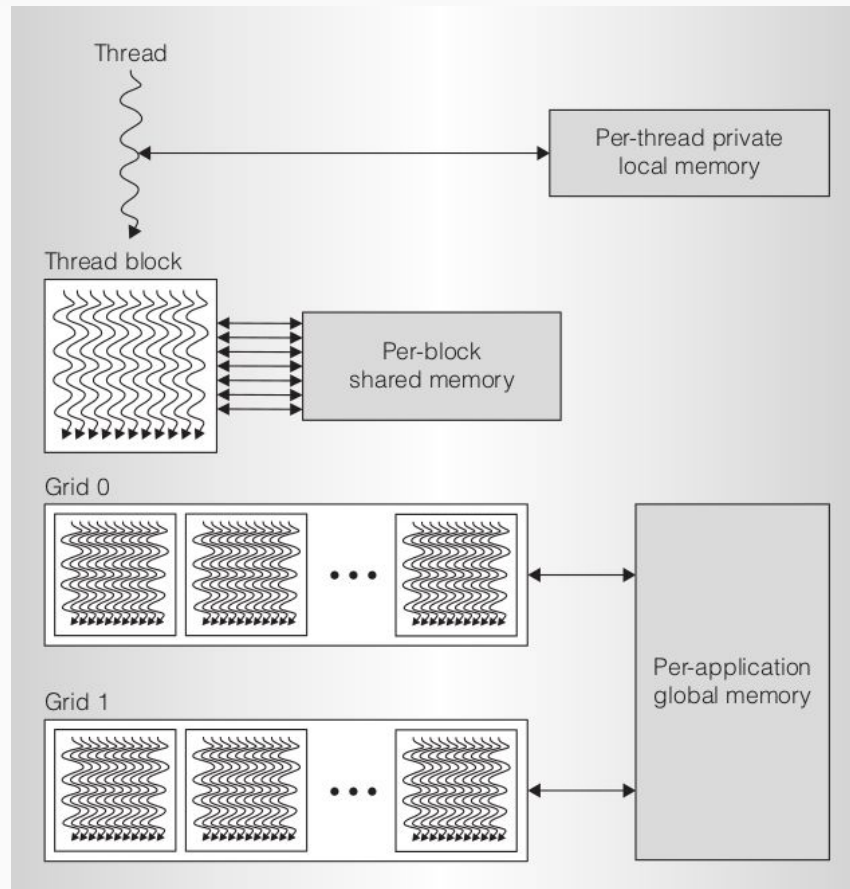


Threads, blocks, and grids

- CUDA arranges kernel execution into

THREADS, **BLOCKS**, and **GRIDS**

- **Thread:**
 - One execution
 - All threads execute the same kernel (but on different data)
 - Has local, private memory

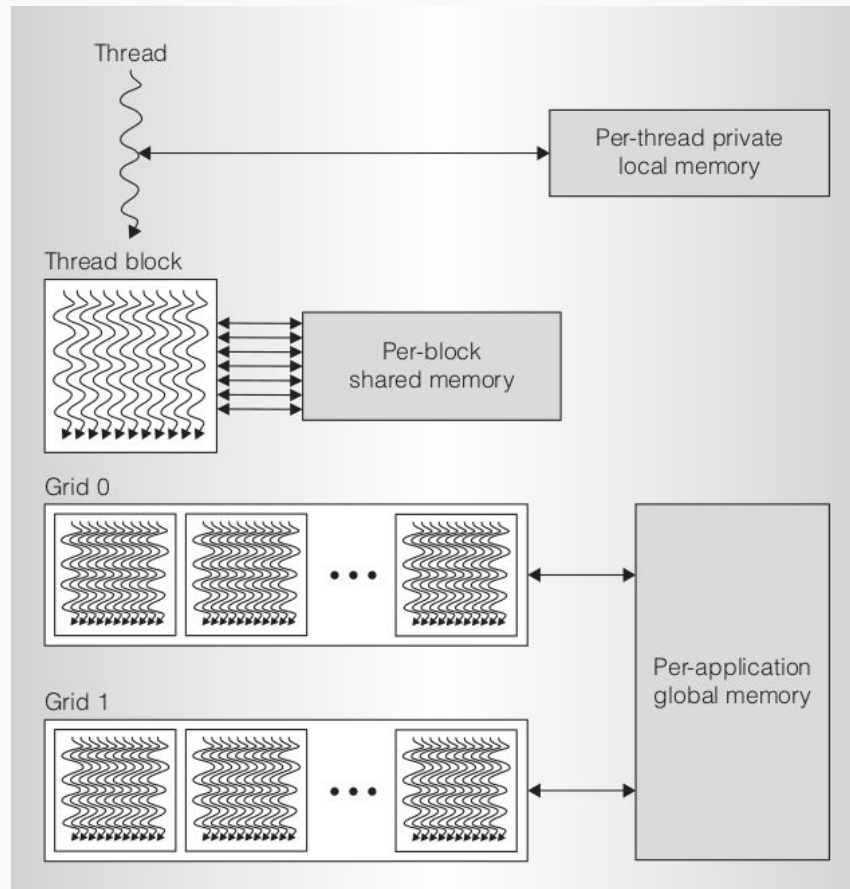


Threads, blocks, and grids

- CUDA arranges kernel execution into

THREADS, **BLOCKS**, and **GRIDS**

- **Thread:**
 - One execution
 - All threads execute the same kernel (but on different data)
 - Has local, private memory
- **Block:**
 - A group of (possibly related) **threads**
 - Has shared memory for all **threads**

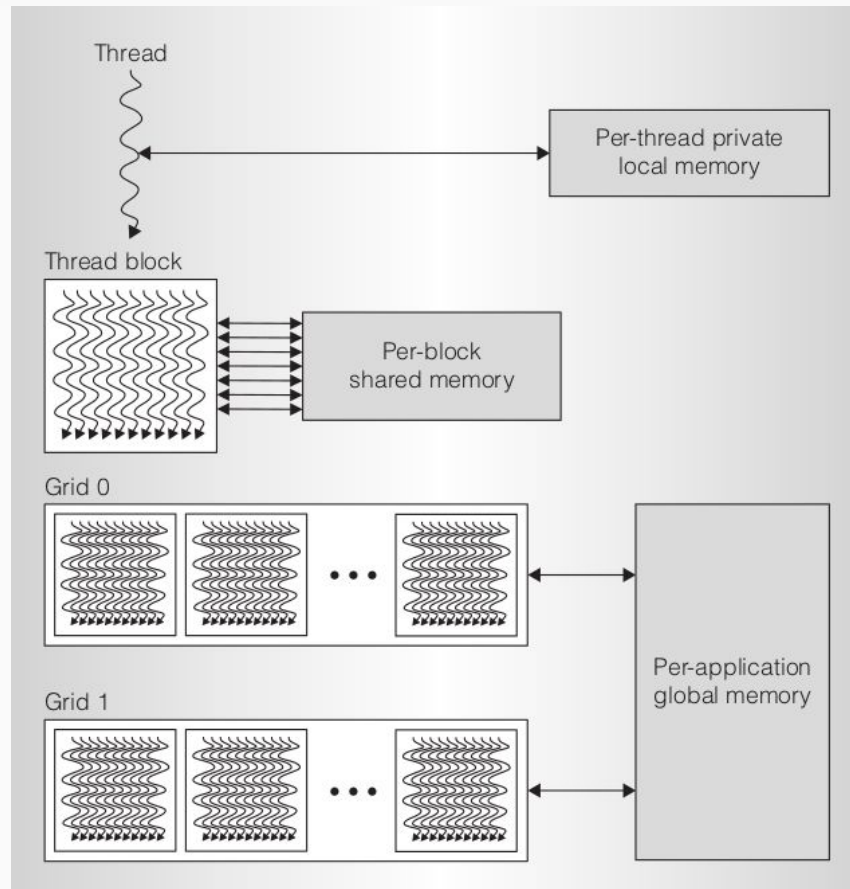


Threads, blocks, and grids

- CUDA arranges kernel execution into

THREADS, **BLOCKS**, and **GRIDS**

- **Thread:**
 - One execution
 - All threads execute the same kernel (but on different data)
 - Has local, private memory
- **Block:**
 - A group of (possibly related) **threads**
 - Has shared memory for all **threads**
- **Grid:**
 - A collection of **blocks**
 - All grid cells have access to shared global memory (application-level)



Example “saxpy”: $y \leftarrow a \cdot x + y$

```
def serial_saxpy(N, a, x, y):  
    for i = 0 .. N-1:  
        y[i] = a * x[i] + y[i]
```

a = <some number>

x = <array of N numbers>

y = <array of N zeros>

serial_saxpy(N, a, x, y)

- Single-precision $ax+y$
- Cost for serial implementation: $O(N)$
- Computational dependencies are trivially parallel

Example “saxpy”: $y \leftarrow a \cdot x + y$

```
def serial_saxpy(N, a, x, y):  
    for i = 0 .. N-1:  
        y[i] = a * x[i] + y[i]
```

```
a = <some number>  
x = <array of N numbers>  
y = <array of N zeros>
```

```
serial_saxpy(N, a, x, y)
```

```
def cuda_saxpy(N, a, x, y):  
    i = blockIdx.x * blockDim.x + threadIdx.x  
    if i < N:  
        y[i] = a * x[i] + y[i]
```

```
a = <some number>  
x = <array of N numbers>  
y = <array of N numbers>
```

```
d_x ← copy x to GPU
```

```
d_y ← copy y to GPU
```

```
cuda_saxpy<<< ⌈N/256⌉, 256 >>>(N, a, d_x, d_y)
```

```
y ← copy d_y from GPU
```

Example “saxpy”: $y \leftarrow a \cdot x + y$

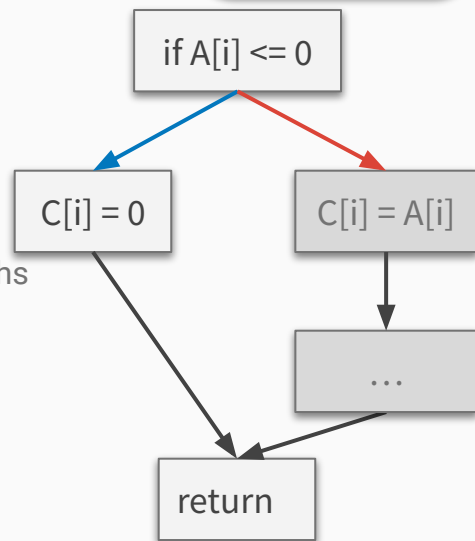
- Data access is managed by the thread:
 - `blockIdx.x` = current block index
 - `blockDim.x` = size of the current block
 - `threadIdx.x` = current thread within block
- $\lceil N/256 \rceil$ = number of **blocks**
256 = number of **threads per block**
- “`if i < N ...`” avoids array bound errors from integer math in block division

```
def cuda_saxpy(N, a, x, y):  
    i = blockIdx.x * blockDim.x + threadIdx.x  
    if i < N:  
        y[i] = a * x[i] + y[i]  
  
a = <some number>  
x = <array of N numbers>  
y = <array of N numbers>  
d_x ← copy x to GPU  
d_y ← copy y to GPU  
  
cuda_saxpy<<<  $\lceil N/256 \rceil$ , 256 >>>(N, a, d_x, d_y)  
  
y ← copy d_y from GPU
```

Thread execution and warps



- **Threads** within a block execute in parallel via SIMT
 - Single-Instruction, Multiple-Thread
- Threads run in groups of 32 called a **warp**
 - Threads start at the same instruction, **but can follow different execution paths**
 - **Warp finishes when all threads have finished**
 - Pre-Volta (2017) GPUs share the instruction counter, disable cores on different paths
 - For maximum efficiency: **threads should follow a common path!**
 - Threads can be explicitly synchronized: `__syncthreads();`
- **Blocks** need not execute all simultaneously
 - E.g. if you have more blocks than cores
 - This is why you can't share memory between blocks



Example 2: Dot product (single-block)

```
// Vector dot product with shared memory
__global__ void dot(float *a, float *b, float *c, int N)
{
    // Each thread acts as a mapper
    __shared__ float temp[N];
    int i = threadIdx.x;
    temp[i] = a[i] * b[i];

    // Wait here for all threads to finish
    __syncthreads();

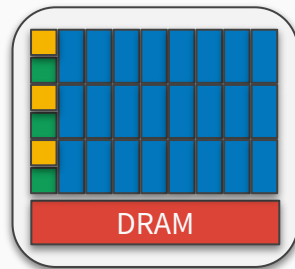
    // Use thread 0 as reducer
    if (threadIdx.x == 0) {
        float sum = 0.0;
        for (int i = 0; i < N; i++) {
            sum += temp[i];
        }
        *c = sum;
    }
}
```

This could be improved in several ways...

- Use blocks to compute partial sums
 - \Rightarrow like **combiners** in map-reduce!
- Divide-and-conquer for sum aggregation
- Use atomic updates for output buffer
 - \Rightarrow like **accumulators** in spark!

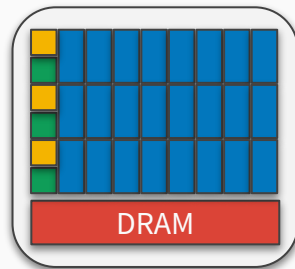
Threads, blocks, grids, and hardware

- **Like mappers**, kernels must be independent of one another and able to execute **in any order**
- **Unlike mappers**, kernels are not pure functions.
 - In fact, they cannot have return values at all!
 - Outputs are written to pre-allocated memory buffers



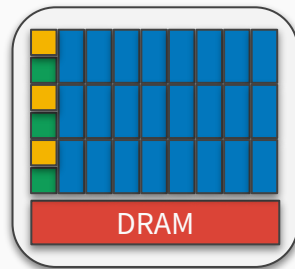
Threads, blocks, grids, and hardware

- **Like mappers**, kernels must be independent of one another and able to execute **in any order**
- **Unlike mappers**, kernels are not pure functions.
 - In fact, they cannot have return values at all!
 - Outputs are written to pre-allocated memory buffers
- To exploit shared memory within blocks, you must be careful about organizing your data (**just like RDD partitions!**)



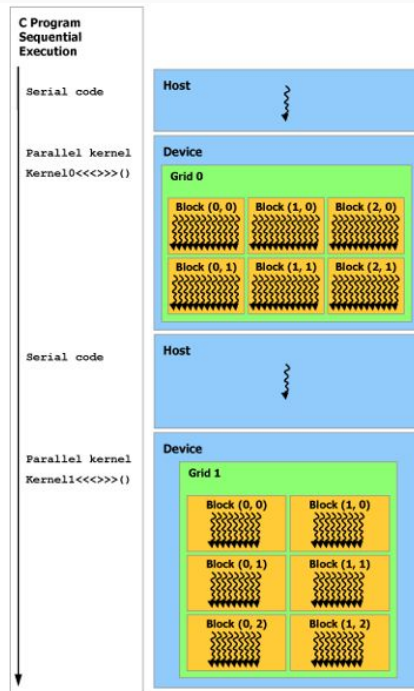
Threads, blocks, grids, and hardware

- **Like mappers**, kernels must be independent of one another and able to execute **in any order**
- **Unlike mappers**, kernels are not pure functions.
 - In fact, they cannot have return values at all!
 - Outputs are written to pre-allocated memory buffers
- To exploit shared memory within blocks, you must be careful about organizing your data (**just like RDD partitions!**)
- Different hardware will have different # cores
 - Thread/block abstraction hides the hardware core layout from the software
 - Current GPUs allow up to 1024 threads per block



Complex programs

- CUDA kernels are usually “simple” operations
 - Vector addition, matrix multiplication, etc.
- Usually we want to combine these to make complex programs
 - E.g. back-propagating through a convolutional network
- Not dissimilar from spark interleaving **transformations** with **actions**



Pitfalls of GPU usage

- Efficient usage relies on keeping the GPUs busy at all times
- Usually an idle GPU is waiting for data from the CPU
 - Memory transfer (**communication**) is often the biggest bottleneck!
 - Try to keep as much data on the GPU as possible without transfers
- Programs with complex control flow (eg conditionals)
- Writing custom GPU code is a daunting task... use existing frameworks!

Summary

Part 2: GPGPUs

- GPGPUs were originally designed to overcome imbalance in graphics, but proved to be a very useful abstraction!
- Shaders \Rightarrow Kernels
- Distributed programs with simple control flow on shared memory