



NYU

Center for  
Data Science

# Week 03.3: Map-Reduce

DS-GA 1004: Big Data

# Announcements

- Quiz #1 on SQL/RDBMS


Friday (02/11)

- You have 24 hours to start the quiz
- You have 1 hour to complete it
- Open-book/paper/note
- **You may not discuss with classmates / other people!**

# Previously...

- Relational data model + DBMS
- Transactions for concurrent access

# This week



1. Introduction to Map-Reduce  
(Dean & Ghemawat, 2008)
2. Criticisms of Map-Reduce  
(DeWitt & Stonebraker, 2008)

# Why map-reduce?

- Distributed programming is really **difficult**!
- If we **restrict** how we program, parallelism becomes easier
- The **map** and **reduce** operations are surprisingly powerful!

# Map-Reduce flow

## 1. Map phase

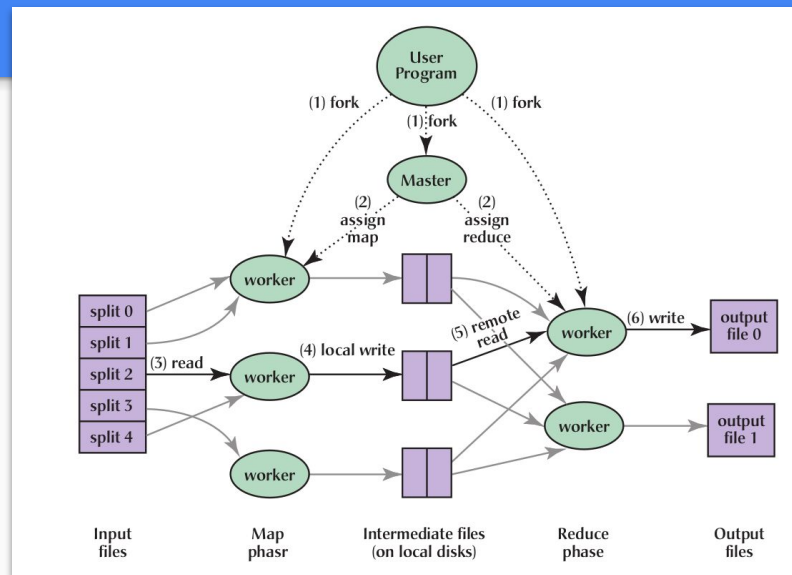
- Distribute data to mappers
- Generate intermediate results (*key*, *value*)

## 2. Sort / shuffle phase

- Assign intermediate results to reducers (by *key*)
- Move data from mappers to reducers

## 3. Reduce phase

- Execute reducers and collect output



# Why “map” and “reduce”?

- These are common operations in **functional programming**
  - E.g.: LISP, ML, Haskell, Scala...
- **map**(function  $f$ , values  $[x_1, x_2, \dots, x_n]$ )  $\rightarrow [f(x_1), f(x_2), \dots, f(x_n)]$ 
  - **map** : function, list  $\rightarrow$  list
- **reduce**(function  $g$ , values  $[x_1, x_2, \dots, x_n]$ )  $\rightarrow g(x_1, \text{reduce}(g, [x_2, \dots, x_n]))$ 
  - **reduce** : function, list  $\rightarrow$  item

# Ex.: word counting in a document collection

```
mapper(doc_id, doc_contents):  
    for word in doc_contents:  
        emit word, 1
```

```
reducer(word, counts):  
    total_count = 0  
    for count in counts:  
        total_count += count  
    emit total_count
```

Key idea:

- Make the mapper as simple as possible
- Let the MR framework route the intermediate results
- Reduce can be simple as well



# Combiner example: word count

```
mapper(doc_id, doc_contents):  
    for word in doc_contents:  
        emit word, 1
```

```
combiner(word, counts):  
    partial_count = 0  
    for count in counts:  
        partial_count += count  
    emit word, partial_count
```

```
reducer(word, counts):  
    total_count = 0  
    for count in counts:  
        total_count += count  
    emit total_count
```

Mapper node

This works because summation is **commutative** and **associative**:

$$A + B = B + A$$

$$A + B + C = (A + B) + C$$

When that happens, you can re-use the **reducer** code as a **combiner**!

# Group exercise 1

Degree counting

- Input consists of weighted edges in an undirected graph, e.g.:

```
# u, v, w  
3, 5, 0.5  
2, 3, 10  
3, 1, 6.1  
...
```

- Write a **mapper** and **reducer** to compute the **degree** (sum of edge weights) for each vertex
  - $\text{deg}(u) = \sum_{u,v} w(u, v) + \sum_{v,u} w(u, v)$

slido



**Write your mapper and reducer code here (one submission per group)**

① Start presenting to display the poll results on this slide.

# Solution: Degree counting

```
def mapper(u, v, w):
```

```
    emit u, w
```

```
    emit v, w
```

```
def reducer(u, weights):
```

```
    emit sum(weights)
```

Each edge touches two vertices, and contributes to both of their degree totals.

⇒ Mapper produces two intermediate keys per edge

# Some tips...

- **Don't use floating point** / real numbers for **keys**
  - Keys need to hash consistently, and floating point equivalence can be subtle!
- Keep **map** and **reduce** simple!
  - Avoid loops if possible
  - Let sort do the work for you
- **Compare** your algorithm's complexity to the **simple / single-core solution**

# Degree-counting, simple algorithm

```
def degree_count(V, E):
```

```
    deg ← defaultdict(0)
```

```
    for (u, v, w) in E:
```

```
        counts[u] += w
```

```
        counts[v] += w
```

```
    return deg
```

- Time is  $O(|E|)$
- Space is  $O(|V|)$

# Degree-counting, MR algorithm

```
def mapper(u, v, w):
```

```
    emit u, w
```

```
    emit v, w
```

```
def reducer(u, weights):
```

```
    emit sum(weights)
```

- **Mapper**

- Each call is  $O(1)$  time/space
- Time  $O(|E| / K)$  for  $K$  mappers
- **Total space:  $O(|E|)$**

- **Shuffle:**  $O(|V|)$  keys,  $O(|E|)$  values

- **Reducer**

- Each call is  $O(\text{max-degree})$
- Total time depends on skew
- Total space is  $O(|V|)$

# Group exercise 2

Directed degree counting

- Input consists of weighted edges in an **directed** graph, e.g.:

```
# u, v, w; u → v  
3, 5, 0.5  
2, 3, 10  
3, 1, 6.1  
...
```

- Write a **mapper** and **reducer** to compute the **in-degree** and **out-degree** for each vertex
  - $\text{in-deg}(u) = \sum_{v,u} w(v, u)$
  - $\text{out-deg}(u) = \sum_{u,v} w(u, v)$



slido



**Write your map-reduce code  
for directed degree counting  
here.**

① Start presenting to display the poll results on this slide.

# Solution: Directed degree counting

```
def mapper(u, v, w):
```

```
    emit (u, 'out'), w
```

```
    emit (v, 'in'), w
```

```
def reducer(key, weights):
```

```
    # u, direction ← key
```

```
    emit sum(weights)
```

The logic is nearly the same as in problem 1

We can pack information into the **intermediate keys** to keep things simple!

Sort/shuffle phase does the hard work for us.

# 5 criticisms

(DeWitt & Stonebraker, 2008)

1. Too low-level
  - Schemas, high-level languages
2. Poor implementation
  - No indexing, “brute-force” thinking
3. Not novel
4. Missing important features
  - Schemas, transactions, etc
5. No DBMS compatibility

slido



**Which of the five criticisms  
do you think are valid?**

① Start presenting to display the poll results on this slide.

slido



**[2] Which of the five criticisms do you think are valid?**

① Start presenting to display the poll results on this slide.

# MR is not an RDBMS...

- ... but there is some overlap!
- Simple queries can often be translated to MR
  - SELECT some\_id, sum(some\_feature) from my\_table  
**WHERE some\_clause** ← mapper → (id, feature)  
**GROUP BY some\_id** ← reducer → (id, sum(feature))
- Complex queries can be very difficult to translate
  - Multi-way joins are especially difficult, and common in practice

# MR is not a general computation engine...

- ... but it can still do a lot!
- Algorithms with **clear parallelism** and no/little looping are okay
- Iterative or recursive algorithms ... not so much
  - Gradient descent (and other ML approaches, e.g. alternating solvers)

# No transactions?

- Transactions exist to maintain consistency and validity of data
- Map-Reduce avoids these issues in a couple of ways:
  - Data is immutable
  - **MR programs must be deterministic**
- If a compute node fails, we can spawn extra instances of mappers / reducers, and we're guaranteed the same outputs
  - As soon as one node finishes, we can ignore any redundant copies



# Next week

- The Hadoop distributed file-system (HDFS)
- Lab 2: programming with map-reduce

slido



# Audience Q&A Session

① Start presenting to display the audience questions on this slide.