



NYU

Center for
Data Science

Week 05.2: Using Spark

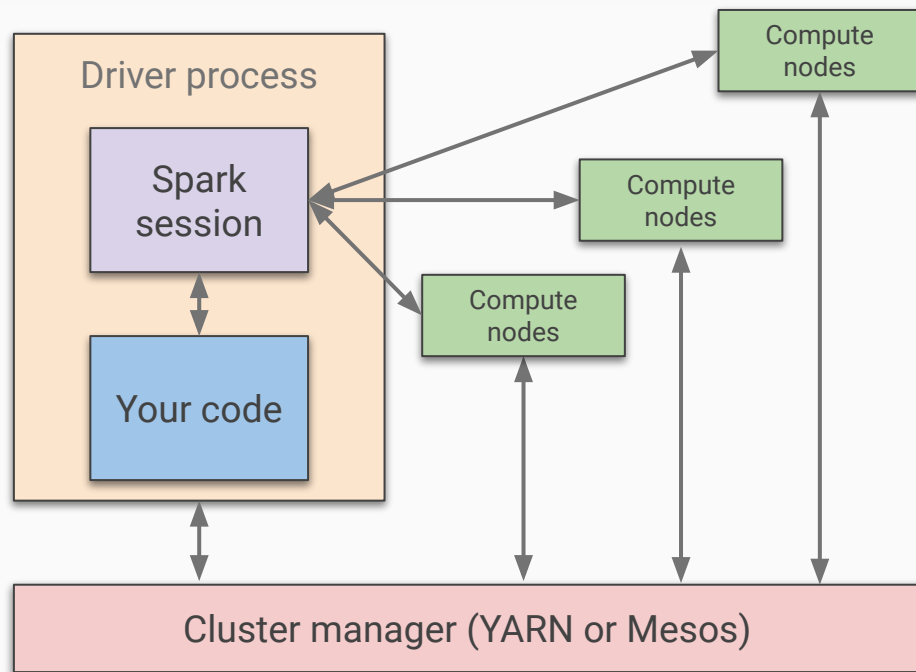
DS-GA 1004: Big Data

Apache Spark (2009, v0 2012)

- Cluster computing framework using RDDs
- Integrates with Hadoop ecosystem
 - HDFS for storage (but other backends are possible)
 - Mesos or (Hadoop) YARN for scheduling
- Written in Scala with API in other languages
 - Python, Java, R, etc

Architecture: session and driver

- **Driver** is the process that you run on, e.g., the head / login node
- The **Session** object connects **your code** to the **cluster** / **compute nodes**



Aside: Why Scala?

- RDD design fits well with functional programming
 - Closures (function + environment) encapsulate everything you need to construct a result
 - Lazy evaluation
 - Immutable data
- Scala compiles to Java virtual machine (JVM)
 - JVM byte code is portable across machines
 - Integration with Hadoop tools (in Java) is relatively easy

Aside part 2: closures

```
function make_closure(x):
```


```
    function f(y):
```

```
        return x + ' stomps on ' + y
```

```
    return f
```

```
stomp ← make_closure('Claosaurus')
```

```
print(stomp('the village'))
```



- **Closures** are a functional programming construction that combine a function with its environment (ie dependencies)
- Does this sound like **RDDs**?
- Example code (\Leftarrow) constructs and then executes a closure (**stomp**)
- Scala's a great language for this!

Example: gradient descent revisited

```
val points = spark.textFile(...).map(parsePoint).cache()
var w = Vector.random(D)

for (i ← 1 to ITERATIONS)
  val grad = spark.accumulator(new Vector(D))

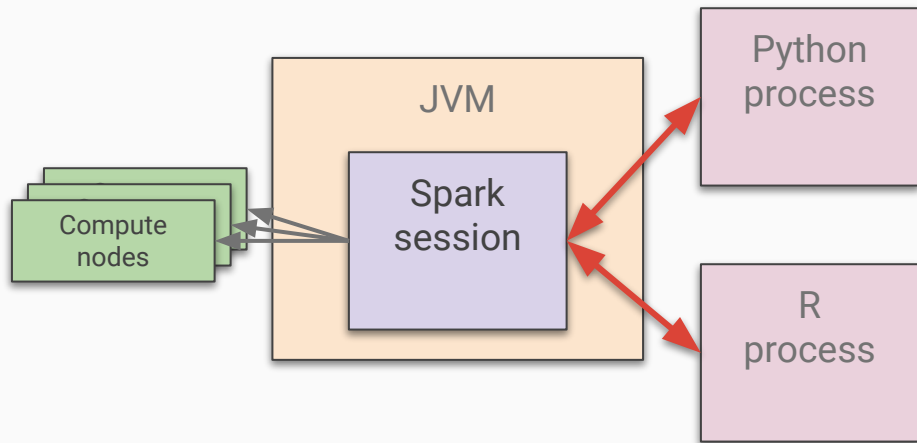
  for (p ← points)
    val grad_p =  $\nabla_w f(p; w)$ 
    grad += grad_p

  w -= grad.value
```

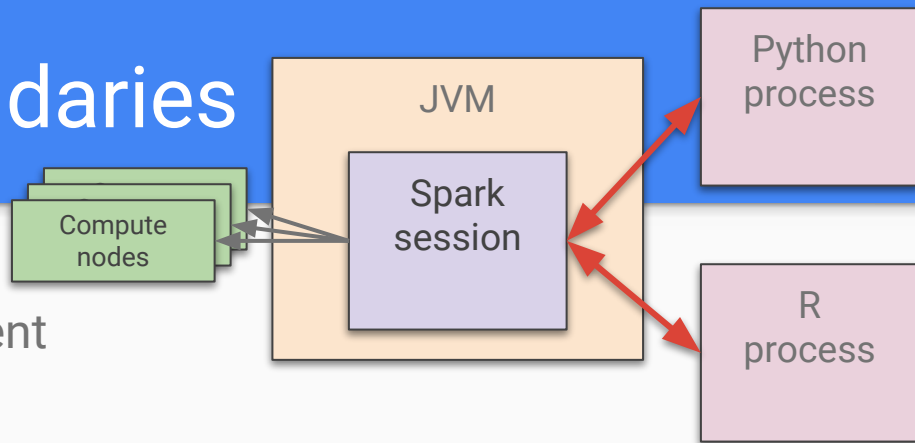
- Some scala notation:
 - **val** means immutable **value**
 - **var** means mutable **variable**
- Outer loop runs in series
- Inner loop runs in parallel over points
 - Equivalent to:
points.foreach(p ⇒ {loop body})
- **grad** is a shared *accumulator*
 - Write-only data structure for associative/commutative updates

Beyond Scala

- You don't need to code Scala to use Spark
- Spark can run from R or Python (or Java)
- Beware: R and Python may not be as fast as Scala
- **Crossing process boundaries** can be expensive, but Spark does a good job of managing this



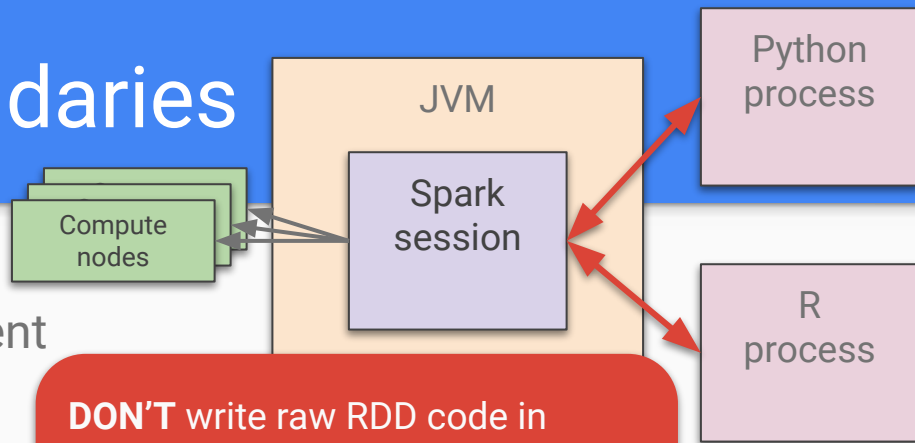
Crossing process boundaries



- Imagine rewriting the gradient descent loop in Python
- Technically it's possible, but it's **slow**
- This is because each operation needs to jump out of **Scala** and into **Python**, serializing all data **between processes**

```
val points =  
  spark.textFile(...).map(parsePoint).cache()  
var w = Vector.random(D)  
  
for (i ← 1 to ITERATIONS)  
  val grad = spark.accumulator(new Vector(D))  
  for (p ← points)  
    val grad_p =  $\nabla_w f(p; w)$   
    grad += grad_p  
  w -= grad.value
```


Crossing process boundaries



- Imagine rewriting the gradient descent loop in Python
- Technically it's possible, but it's **slow**
- This is because each operation needs to jump out of **Scala** and into **Python**, serializing all data **between processes**

DON'T write raw RDD code in Python (or R)

DO use existing packages written in Scala with Python bindings

```
for (i <- 1 to ITERATIONS)
  val grad = spark.accumulator(new Vector(D))
  for (p <- points)
    val grad_p =  $\nabla_w f(p; w)$ 
    grad += grad_p
  w -= grad.value
```

Spark DataFrames API

- RDDs are great, but a bit cumbersome for ad-hoc computation
- **DataFrames** are common representations in many languages
 - R, pandas (Python), etc.
- Spark 2.x added a DataFrame API as a primary interface
 - Code looks more or less like pandas/Python!

DataFrames and RDDs

- DataFrames in Spark are like relations in RDBMS
 - Well-defined schema with types over columns
 - Each row is a tuple (sort of...)
- DataFrames operations are translated into RDD transformations by Spark
- RDD transformations can then be executed within JVM
 - No more serialization of data between JVM \leftrightarrow Python!

Spark-SQL

- Spark 2.x allows you to express queries in SQL
 - Or using an object-method chaining API – the two are equivalent!
- Queries are executed against DataFrames
 - DataFrames are secretly RDDs, not RDBMS tables!
- Queries can be optimized by analyzing the RDD lineage graph

```
df.createOrReplaceTempView('my_table')


res = spark.sql(' SELECT zip_code, sum(height) as H
                  FROM my_table
                  GROUP BY zip_code')

res.show()

df.groupBy('zip_code')
  .sum('height')
  .withColumnRenamed('sum(height)', 'H')
  .show()
```

Repartitioning

- Sometimes you know in advance which columns of a DataFrame will be filtered
 - E.g., dates or timestamps
- You can give hints to Spark that RDD partitions should align accordingly
 - `df.repartition(# PARTITIONS, col("NAME OF COLUMN"))`
 - This can reduce the width of RDD dependencies
- **This is much like indexing in RDBMS**



Tips and pitfalls

- Before running an action, run the `explain()` method on the DataFrame
 - This will give you an execution plan
 - You might identify some inefficiencies or bugs this way
- Be careful with `collect()`!
 - This will stream all results back to the driver node
 - If it's a large data set, and you forgot an aggregation step, this will be very bad news.
 - Test-drive a large query with `take(10)` instead of `collect()`
 - Probably you want `.save()` instead of `.collect()` anyway

Wrap-up on Spark

- RDD framework is more flexible than Map-Reduce
- Caching can make interactive jobs faster
- SparkSQL / DataFrames API makes development easy